

From Bottlenecks to Breakthroughs: Optimizing zkEVM Provers

Leo Jeong
Linea, Consensys

Devcon 7
Nov 13, 2024



- **Intro**
 - **Journey of Optimizing zkEVM Provers in Production**
 - **Why is zkEVM prover efficiency important?**
 - **Characteristics of zkEVM provers and common bottlenecks**
 - **Scope of the talk: what will and won't be covered**
- **Optimizing zkEVM Provers**
 - **Diagnostics and Bottleneck Analysis**
 - **CPU Optimization**
 - **Memory Optimization**
 - **Others**
- **Conclusion**
 - **Results**
 - **Lessons learned: Dos and Don'ts**
 - **Future work**

Intro Journey of Optimizing zkEVM Provers in Production

Initial challenges and problems

- Linea mainnet alpha launch (Aug 16th, 2023)
- Prover was slow
- We were spending \$\$\$ for Provers



Mainnet Alpha
is Here



Intro Journey of Optimizing zkEVM Provers in Production

What we did

- 1 year of optimization in production
- Reduced the prover runtime by 75% (40 to 10 mins)



Intro Why is zkEVM prover efficiency important?

Cost

- One of the biggest cost contributors
- Prover optimization → directly reduces operational expenses

Performance Impact (prover runtime)

- Faster verification and finalization
- Improved throughput of zk-rollups (e.g., TPS).
- Better user experience

Scalability

- Supports more transactions
- Handles more complex proofs

Intro Characteristic of zkEVM Provers and common bottlenecks

Complexity

- Inherently complex due to the full compatibility required with EVM
- Each zkEVM project has its own design

High resource consumption

- CPU and Memory-intensive due to large-scale cryptographic calculations and extensive data structures

Common bottlenecks and issues

CPU usage / parallelization

Memory overhead

I/O bottlenecks

Intro Scope of the talk: what will and won't be covered

What won't be covered

- Linea-specific optimizations
- Detailed cryptographic theory behind zkEVM prover
- Hardware accelerated optimizations (GPU, FPGA, etc.)

What will be covered

- Quick wins and strategies for any zkEVM project
- CPU/memory optimizations without architectural changes or significant modifications

- **Intro**
 - **Journey of Optimizing zkEVM Provers in Production**
 - **Why is zkEVM prover efficiency important?**
 - **Characteristics of zkEVM provers and common bottlenecks**
 - **Scope of the talk: what will and won't be covered**
- **Optimizing zkEVM Provers**
 - **Diagnostics and Bottleneck Analysis**
 - **CPU Optimization**
 - **Memory Optimization**
 - **Others**
- **Conclusion**
 - **Results**
 - **Lessons learned: Dos and Don'ts**
 - **Future work**

Optimizing zkEVM Provers: Diagnostics and Bottleneck Analysis

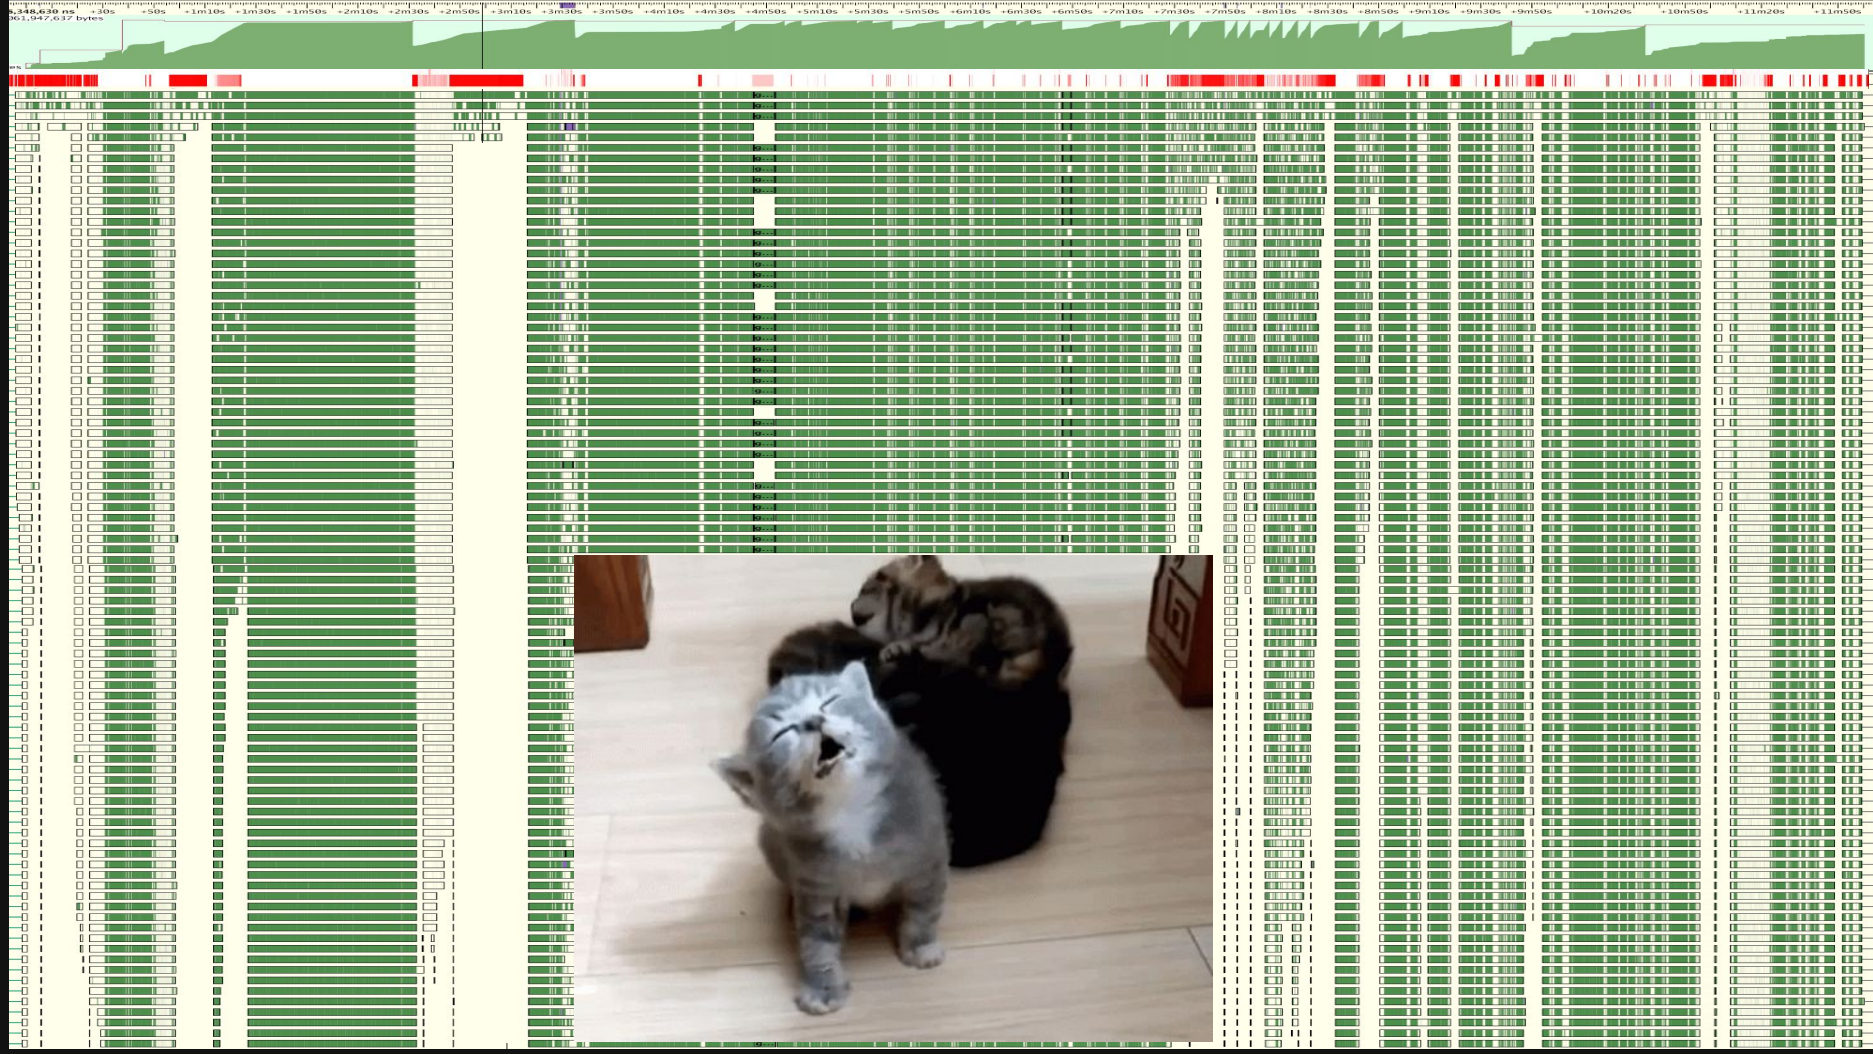
Measure before you optimize


- "Premature Optimization Is the Root of All Evil"
by Donald Knuth, The Art of Computer Programming

CPU and memory profiling

- Visualize CPU and memory usage
- Understand your code
- Built-in or 3rd-party profiling tools
 - go lang: <https://gotraceui.dev/>



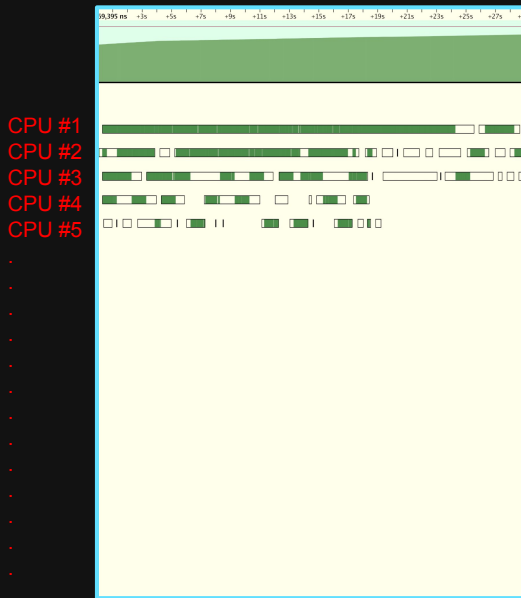




Optimizing zkEVM Provers: Diagnostics and Bottleneck Analysis

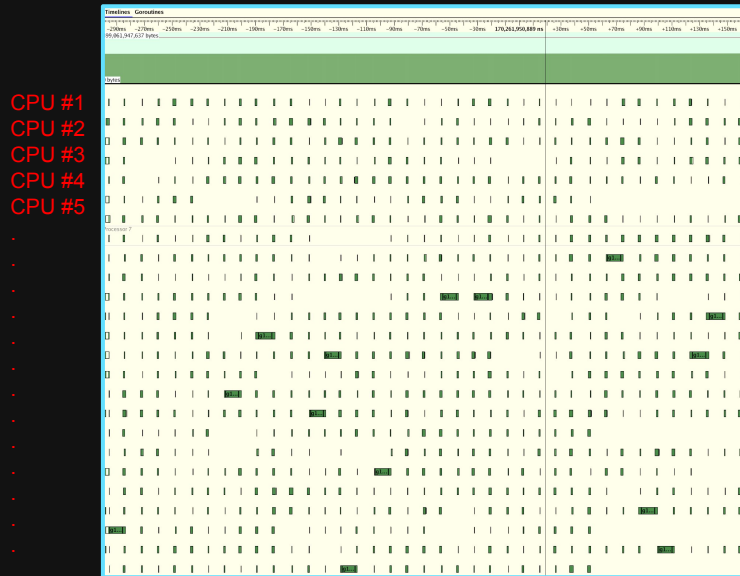
Poor Parallelization:

- Case 1, 3, 4, 6 7 and 8





CPU Underutilization:

- Case 2, 5, 7 and 8

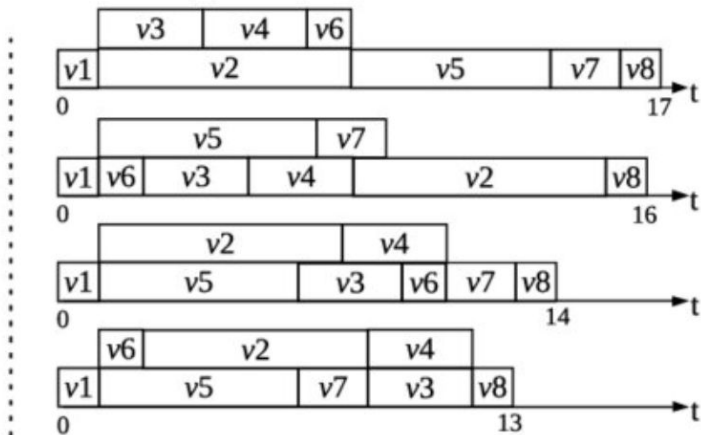


Optimizing zkEVM Provers: CPU Optimization

Leveraging DAG (Directed Acyclic Graph) for Better Parallelization



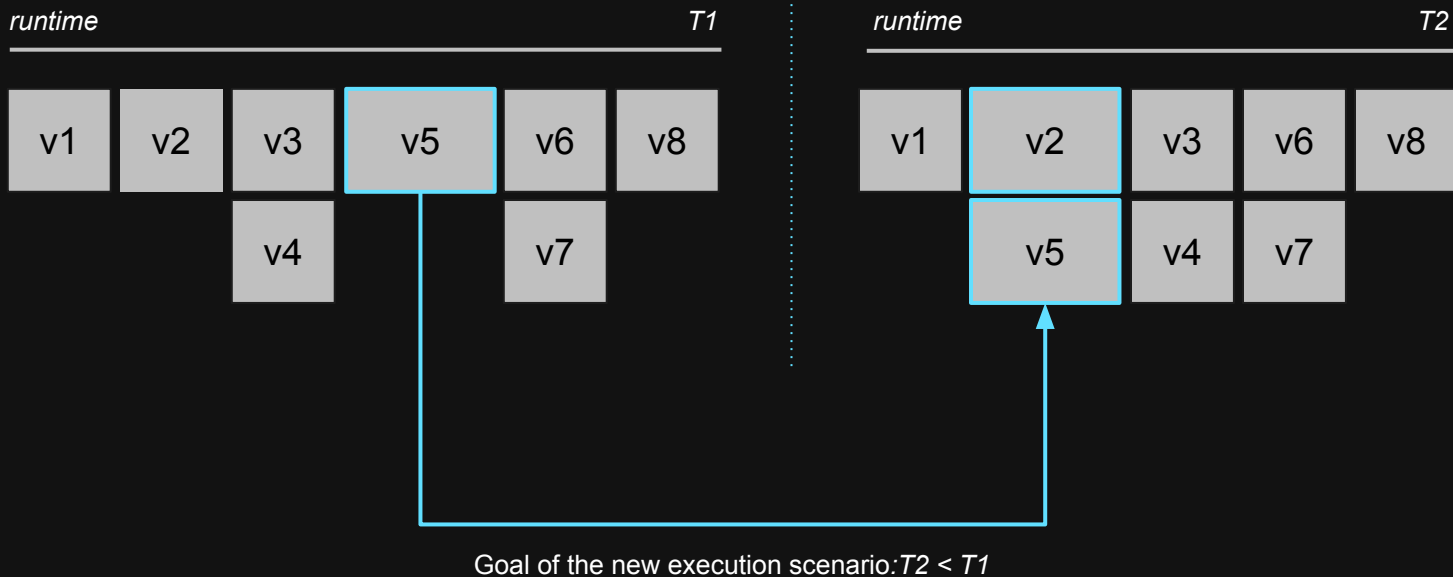

(a) Example DAG



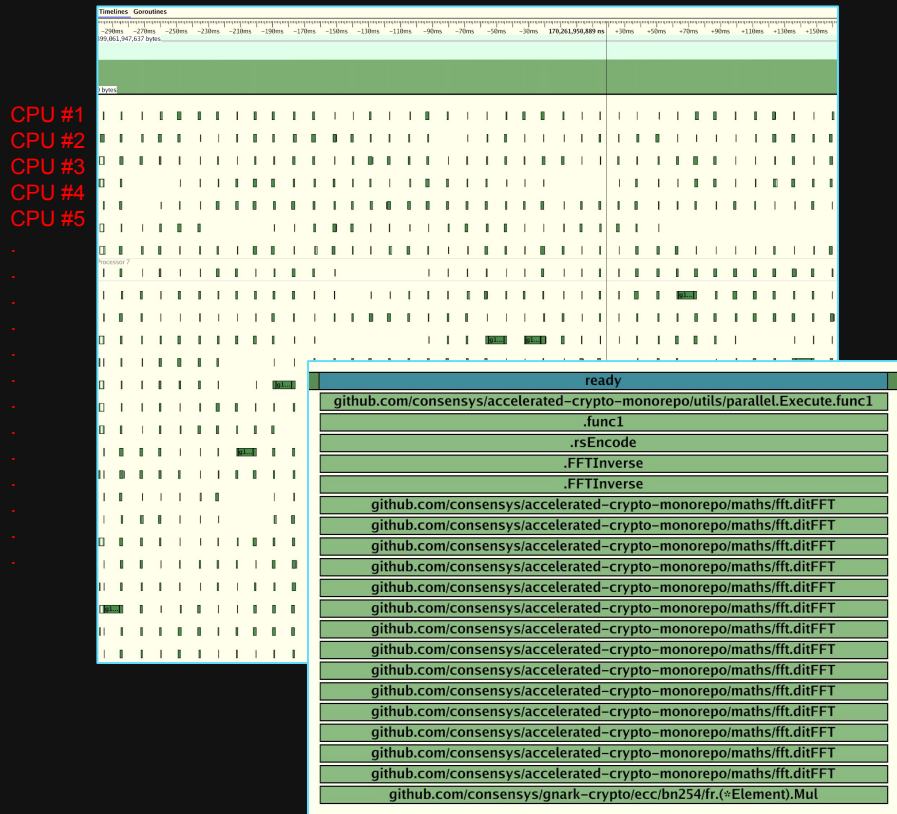
(b) Execution scenarios

Optimizing zkEVM Provers: CPU Optimization

Leveraging DAG (Directed Acyclic Graph) for Better Parallelization



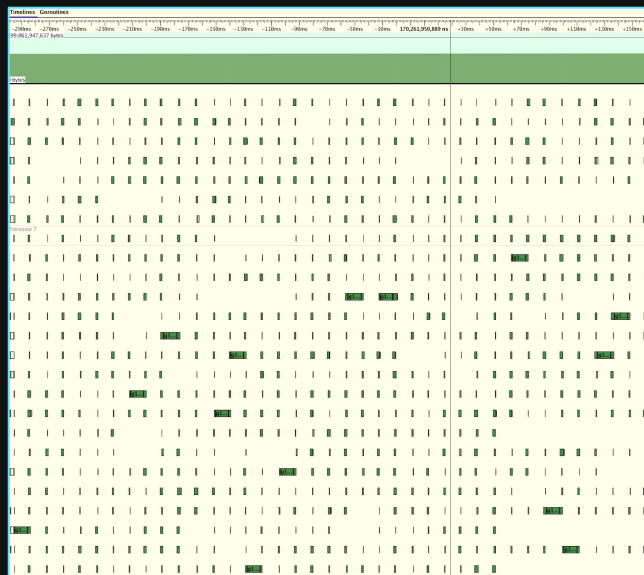
- Poor parallelization
- Synchronization
- Locks
- Blocking I/O operations (e.g., file reads, network requests)




Optimizing zkEVM Provers: CPU Optimization




Poor Parallelization



Underutilization



CPU #1
CPU #2
CPU #3
CPU #4
CPU #5



- **Optimizing zkEVM Provers**
 - **Diagnostics and Bottleneck Analysis**
 - CPU and Memory Profiling
 - Identifying bottlenecks
 - **CPU Optimization**
 - Poor Parallelization (DAG)
 - Underutilization
 - **Memory Optimization**
 - **Others**



Optimizing zkEVM Provers: Memory Optimization [peak: 700 GB]

Why is it important?

- Memory is expensive
 - hpc6id.32xlarge (64 cores, 970 GB): 5.70 USD / Hour
 - hpc6a.xlarge (96 cores, 380 GB): 2.88 USD / Hour
- Higher peak memory usage → more expensive machines

Goal


- Flatten memory usage and eliminate memory peaks



Optimizing zkEVM Provers: Memory Optimization [700 → 450 GB]

Memory pool and allocation

- Define a large memory pool of frequently used data structure
- Reduce the number of memory allocation
- Reduce GC



```
type FromSyncPool struct {
    size int
    P     sync.Pool
}

func CreateFromSyncPool(size int) *FromSyncPool {
    return &FromSyncPool{
        size: size,
        P: sync.Pool{
            New: func() any {
                res := make([]field.Element, size)
                return &res
            },
        },
    }
}

func (p *FromSyncPool) Alloc() *[]field.Element {
    res := p.P.Get().(*[]field.Element)
    return res
}

func (p *FromSyncPool) Free(vec *[]field.Element) error {
    if len(*vec) != p.size {
        utils.Panic("expected size %v, expected %v", len(*vec), p.Size())
    }

    p.P.Put(vec)



    return nil
}
```

Optimizing zkEVM Provers: Memory Optimization [450 -> 370 GB]

Memlimit

- Soft-limit of memory peak by triggering GC when the memory usage hits MEMLIMIT (*MEMLIMIT=370GB*)
- ~~hpc6id.32xlarge (64 cores, 970 GB): 5.70 USD / Hour~~
- hpc6a.xlarge (**96 cores, 380 GB**): 2.88 USD / Hour

MEMLIMIT=370GB



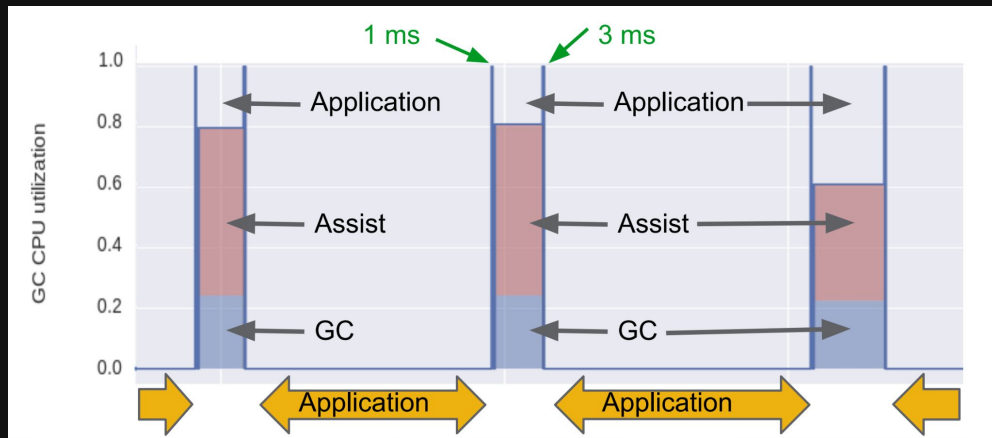
Optimizing zkEVM Provers: Memory Optimization [370 GB]

Garbage-collected language

- Go, Java, Python, ...

It's convenient, but...

- Non-deterministic GC is a problem
- What if it happens during a critical task?



reference: <https://go.dev/>

Memory Optimization [370 GB] + Manual GC

Fine-grained garbage collection

- Make sure no GC is triggered during memory and CPU-intensive tasks
- Manage non-deterministic GCs for critical areas of your code

Manual GC

- Run manual GC before critical tasks
- Less context switching overhead
→ less GC pause time
- Significant runtime improvement

```
if ctx.DomainSize >= GC_DOMAIN_SIZE {
    runtime.GC()
}

go func() {
    // Compute once the FFT of the natural columns
    parallel.ExecuteChunky(len(ctx.AllInvolvedRoots), func(start, stop int) {
        for k := start; k < stop; k++ {
            pol := ctx.AllInvolvedRoots[k]
            name := pol.GetColID()

            var witness sv.SmartVector
            lockRun.Lock()
            witness, isNatural := run.Columns.TryGet(name)
            lockRun.Unlock()

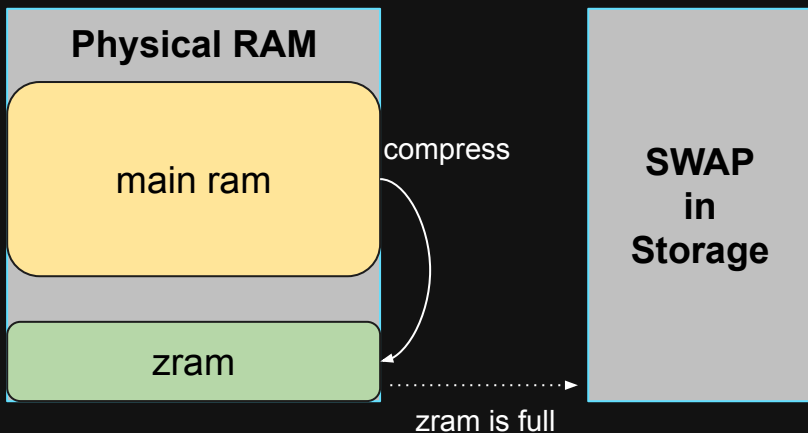
            if !isNatural {
                witness = pol.GetColAssignment(run)
            }
            witness = sv.FFTInverse(witness, fft.DIF, false, 0, 0, nil)

            lock.Lock()
            coeffs[name] = witness
            lock.Unlock()
        }
    })
    wg.Done()
}()
```

Optimizing zkEVM Provers: Memory Optimization [No OOM]

ZRAM (In-memory swap)

- To handle extensive memory demands and avoid Out-of-memory (OOM)
- Available since linux kernel version 3.14
- Massive integration tests with TBs of memory usage
 - `vm.swappiness = 1 to 10`
 - `lz4 compression algorithm` (<https://github.com/lz4/lz4>)



750 GB physical memory + 1.5 TB ZRAM In-memory Swap Space

- **Optimizing zkEVM Provers**

- **Diagnostics and Bottleneck Analysis**

- CPU and Memory Profiling
 - Identifying bottlenecks

- **CPU Optimization**

- Poor Parallelization (DAG)
 - Underutilization



- **Memory Optimization**

- Memory Pooling
 - Memlimit
 - Fine-grained Garbage Collection (Manual GC)
 - ZRAM



- **Others**

Optimizing zkEVM Provers: Others


Data Compression

- The transaction traces
→ show high compression ratio
- 40GB → 2GB
- Use it when you have I/O bottlenecks



Zeroization

- Time taken for initializing your data structure
- Auto-zeroization, optimized libraries for zeroization, etc.



- **Intro**
 - **Journey of Optimizing zkEVM Provers in Production**
 - **Why is zkEVM efficiency important?**
 - **Characteristics of zkEVM provers and common bottlenecks**
 - **Scope of the talk: what will and won't be covered**
- **Optimizing zkEVM Provers**
 - **Diagnostics and Bottleneck Analysis**
 - **CPU Optimization**
 - **Memory Optimization**
 - **Others**
- **Conclusion**
 - **Results**
 - **Lessons learned: Dos and Don'ts**
 - **Future work**

Conclusion

Results: Happy, Happy, Happy

- Reduced the zkEVM prover runtime by 75%
- Saved operating cost
- Improved Prover throughput
 - Faster finality
 - Higher TPS
- (Ironically) We understood what we did and what we are doing



Conclusion

Lessons learned: Dos and Don'ts for other zkEVM projects

- Don't pre-optimize your code (premature optimization)
- Make optimization plans based on measurement
- Do it with a global view to set priorities

Future work

- Distributed zkEVM prover
 - different modules on different machines
- Hardware accelerated zkEVM prover

Optimization is never a one-time task.

Measure, Analyze, Optimize and Repeat.



<https://github.com/Consensys/linea-monorepo>

THANK YOU

leo.jeong@consensys.net