

JSONRPC Enhancement in Geth

@jsvisa
AmberGroup



The Project proposal

The Project Proposal(background)

Currently:

- Use debug_traceXXX RPC to get the block traces
- While debug_traceXXX based on the historical statedb to replay the transactions
 - Slow performance
 - High CPU usage

The Project Proposal

Introduce a real-time tracing system, which will

- offer more efficient and prompt insights into the EVM execution
- support indexing the live execution data

How to index the live data

Workflow/Indexing

Support indexing:

- Traces: block number -> traces
- Nonce: (tx sender's address + nonce) -> Tx Hash

Workflow/Data Schema

kvdb: used to store the newly created block traces and nonces

- A generic Key-Value Database(ethdb)
- Key Schema:
 - Traces: ``BlkNum || BlkHash || TraceType``
 - Nonce: ``Tx Sender's address || nonce``
- Value Schema: RLP encoded data

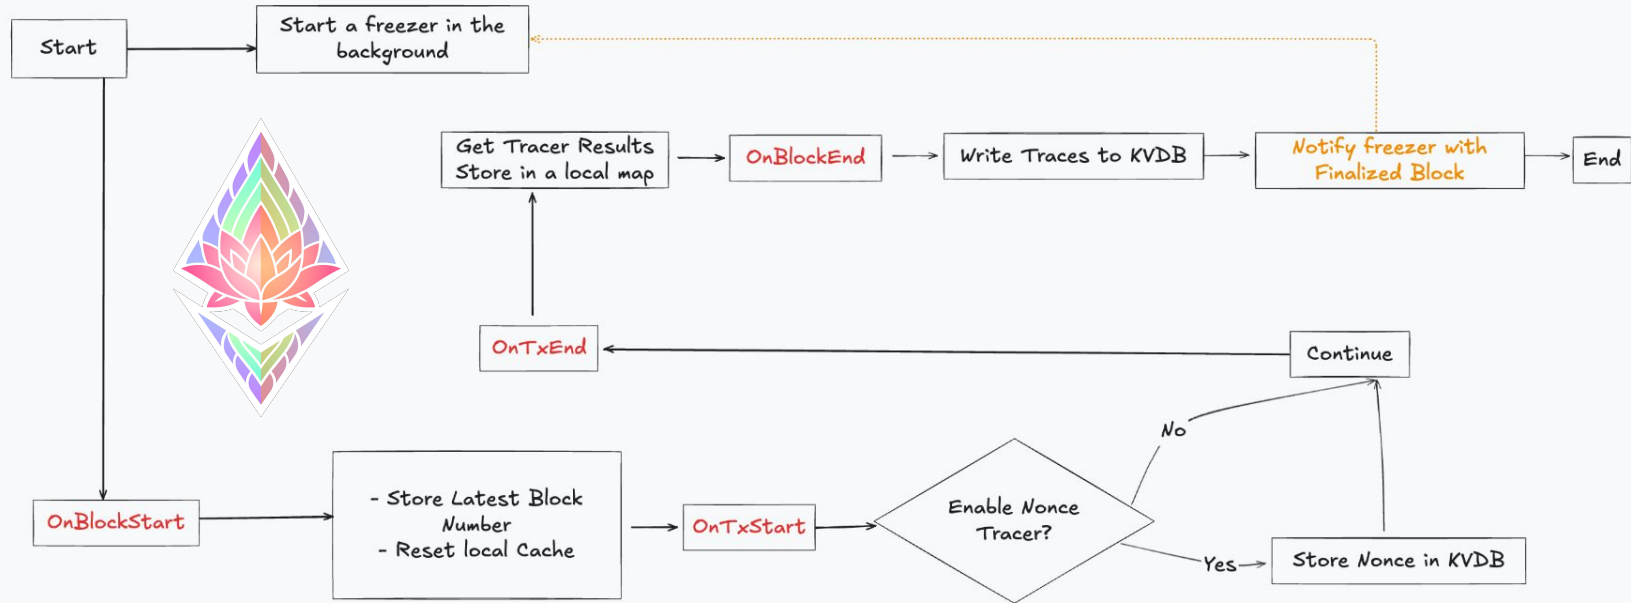
Workflow/Data Schema

kvdb: used to store the newly created block traces and nonces

- A generic Key-Value Database(ethdb)
- Key Schema:
 - Traces: ``BlkNum || BlkHash || TraceType``
 - Nonce: ``Tx Sender's address || nonce``
- Value Schema: RLP encoded data
- TraceType is one of:
 - *callTracer*
 - *flatCallTracer*
 - *parityTracer*
 - *prestateTracer*

Workflow/indexing

Indexing



How to retrieve the data

Workflow/Retrieving

Traces:

- Retrieve by block number/hash
- Retrieve by transaction hash
 - Fetch block number by transaction first -> block number

Nonce:

- Retrieve by sender's address + nonce

Workflow/Retrieving block traces

Retrieve Process





How to use this feature

How to use this feature/setup

Haven't merged yet, see [#30255](#) for more detail

```
$ geth ... --http.api=...,trace \
  --vmtrace=live \
  --vmtrace.jsonconfig='{
    "path": "/data/live-trace",
    "enableNonceTracer": true,
    "maxKeepBlocks": 100000,
    "config": {
      "callTracer": { "withLog": true },
      "parityTracer": {},
      "prestateTracer": { "diffMode": true }
    }
  }'
```

You need to enable the trace namespace either in [http.api](#) or [ws.api](#)

Configuration:

- path: a directory used to store the indexed data
- enableNonceTracer: whether or not to indexing the tx sender's tx-hash
- maxKeepBlocks: keep at most N blocks, older ones will be pruned periodically
- config: a map of each tracers you want to trace with, see [builtin-tracers](#) for more detail

How to use this feature/retrieve

Supported RPCs:

- trace_block
- trace_filter
- trace_transaction
- eth_getTransactionBySenderAndNonce

How to use this feature/retrieve

Supported RPCs:

- trace_block
- trace_filter
- trace_transaction
- eth_getTransactionBySenderAndNonce

- The trace_* RPCs are compatible with [Parity's](#)
- Besides of the Parity's call tracer, we can also retrieve other geth's native tracers' results.

How to use this feature/retrieve

Supported RPCs:

- trace_block
- trace_filter
- trace_transaction
- eth_getTransactionBySenderAndNonce
- The trace_* RPCs are compatible with [Parity's](#)
- Besides of the Parity's call tracer, we can also retrieve other geth's native tracers' results.

Eg: get the block(0x1)'s prestateTracer results

```
{
  "id": "1",
  "jsonrpc": "2.0",
  "method": "trace_block",
  "params": [
    "0x1",
    {"tracer": "prestateTracer"}
  ]
}
```

The performance



Space-Time trade-off

The performance between trace_ and debug_trace

Pro:

- No need to re-apply based on a state
- $O(1)$ time complexity to retrieve the data
- Realtime retrieving

Con:


- Need extra disk space
- Little impact on the on-chain block syncing

Let's compare the results

Use etl to dump traces from geth RPC into Postgres

```
$ ./etl dump2 \  
  --chain=ethereum \  
  --lag=64 \  
  --start-block=21010865 \  
  --end-block=21054092 \  
  --period-seconds=60 \  
  --max-workers=10 \  
  --block-batch-size=10 \  
  --batch-size=1 \  
  ...
```


Export from JSON-RPC latency(ms)



Thank you!

@jsvisa
AmberGroup

