

Leveraging High-Performance computing for Efficient STARK Provers

Ricard Borrell Pol
Staff Engineer, Polygon



Outline

- STARK Implementation in Polygon zkEVM
 - Understanding Bottlenecks in STARK Proof Generation
 - Optimizations
- New STARK prover for ZISK zkVM
 - Divide and conquer: VADCOPs
 - Distributed Prover Architecture

STARK implementation in Polygon zkEVM

STARK in Polygon zkEVM

- STARK with support for subset of arguments
- Prime: Goldilocks
- Number of polynomials 1335
- Degree bound of polynomials 2^{25}
- Main memory requirements: 850 Gb
- Recursion for compression and aggregation

Loop stages

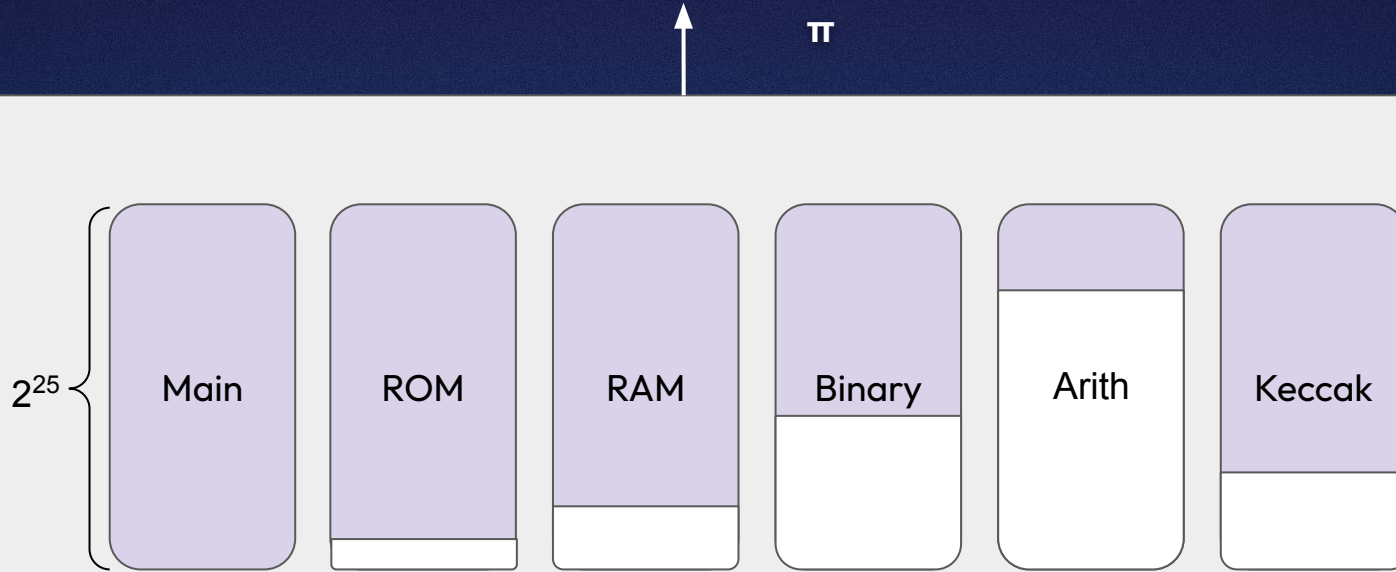
- Calculate challenge
- Calculate witness
- Commit (LDE+Merklee Tree)

**commit
phase**

FRI proximity proof

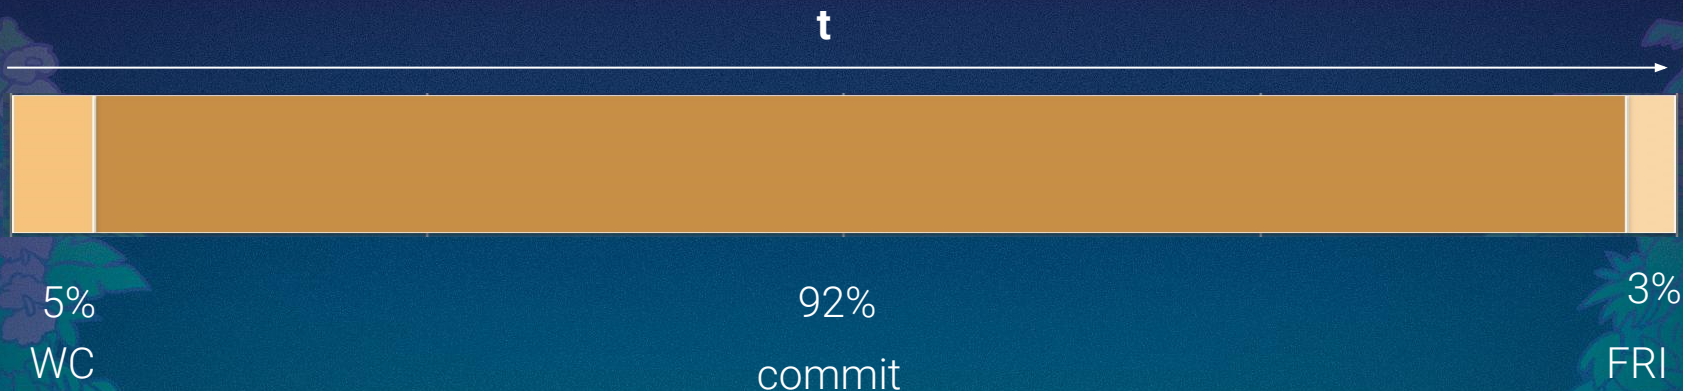
**openings
phase**

STARK in Polygon zkEVM



Current Proof: Monolithic

STARK implementation at Polygon zkEVM



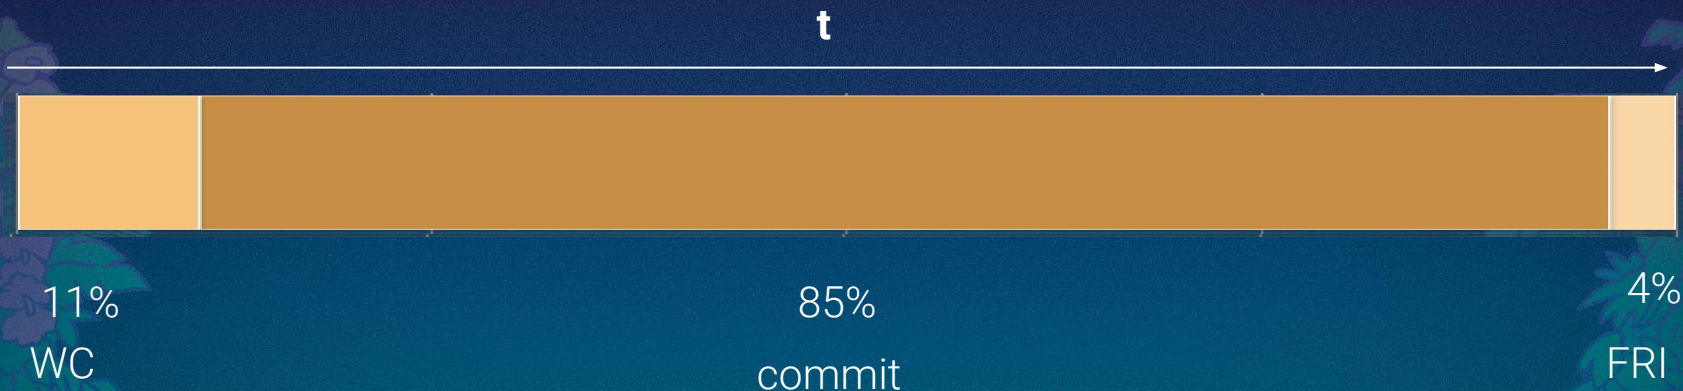
Total time: 359s

Prove cost: 0.46\$

Cost ETH transactions ~0.0002

Cloud instance c3-highmem-176: 176 vCPU (88 Cores - AVX512) y 1.4TB de RAM... ~4.64\$/hour (spot)

STARK implementation at Polygon zkEVM




Total time: 191s

Prove cost: 0.26\$

Cost ETH transactions ~0.0001

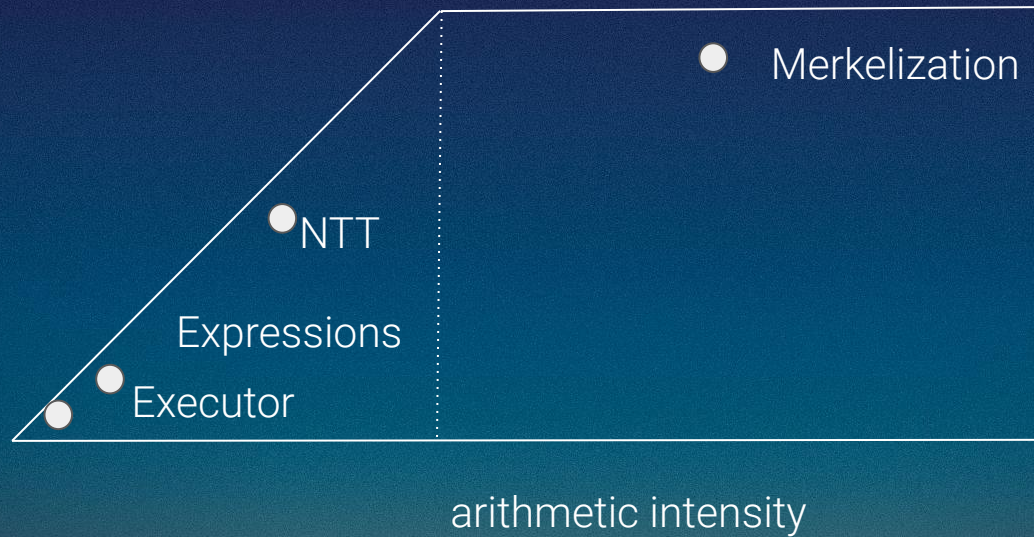
Hardware: Intel(R) Xeon(R) Gold 6462C, 64 core 128, 3.9 GHz, 4x GeForce RTX 4090, 1TB (4.84\$/hour)

STARK implementation at Polygon zkEVM



NTT + MT + EXPRESSIONS + EXECUTOR = 95% prove time




Roofline Model



```
if(omp_get_num_thread()==0){  
}
```



Optimizations: NTT

- Number Theoretic Transform (43%)
- Focus only on minimizing L3 cache misses
- Two key ideas:
 - Parallel FFT by blocks
 - Compute several columns at once



Optimizations: Merkle Tree

- Poseidon hash
- Main improvements obtained from AVX and AVX512





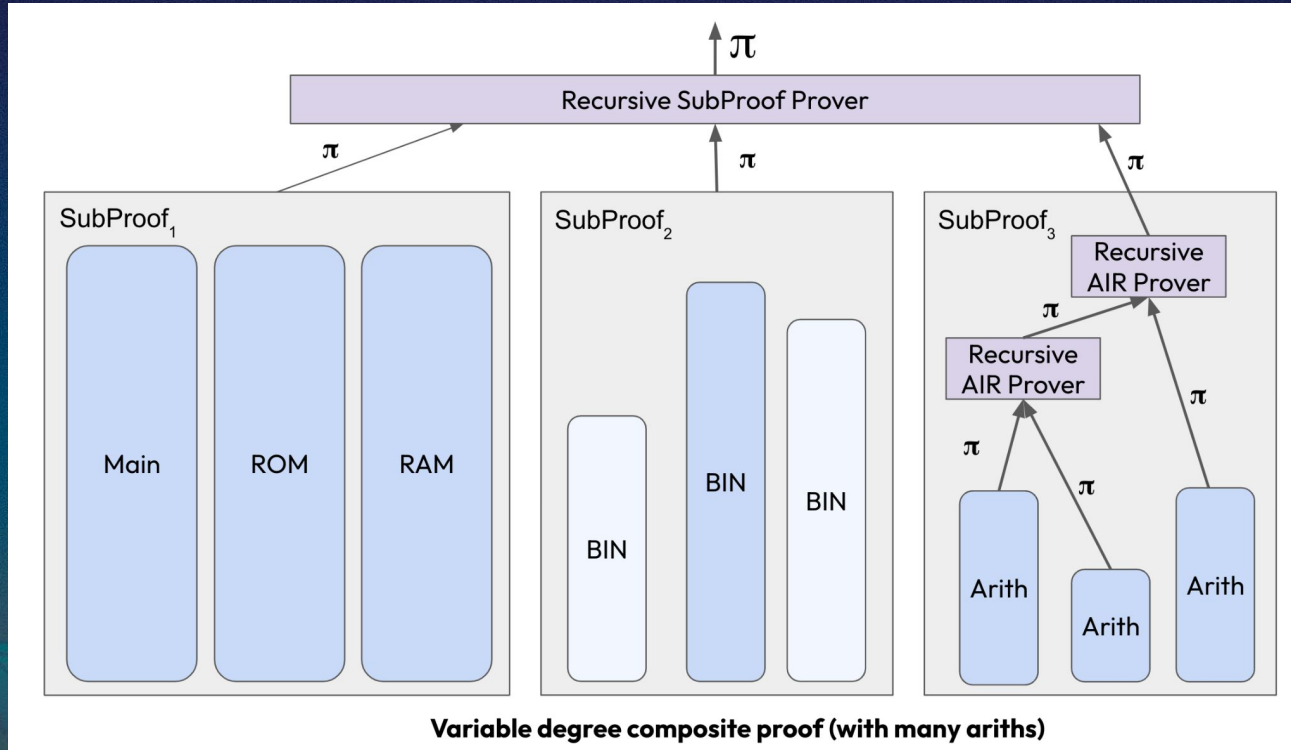
The zkEVM prover fully meets its application needs

New STARK prover for the ZISK zkVM


zisk zkVM

- **Zero Knowledge Virtual Machine** (zkVM)
- Prove execution of arbitrary program written in Rust (or any other high level language)
- Motivation: leverage all the knowledge, experience and tooling gained from the development and maintenance of zkEVM mainnet
- Open repository: <https://github.com/0xPolygonHermes/zisk>

Variable Degree Composite Proofs



Recursive SubProof Prover



subproofs

compressor

recursive 1

recursive 2

Computational Implications

1. CPU efficiency does not improve, moreover ,we add aggregation costs!
2. But new air instance granularity unlocks many optimizations
 - **Accelerators** can be used without PCIe bottlenecks
 - **Vectorized STARK**: Concurrent resolution of multiple instances on vector registers
 - **Distributed prover**: distribute computation across multiple compute nodes

Distributed Prover

Sellenius supercomputer

- HPC resources were granted for public **research project** about STARKS scalability on supercomputers
- **CPU nodes:** 2 AMD Rome or Genoa CPUs
2.4GHz 128 cores and 256 GiB. Up to 208,896 CPU cores
- **GPU nodes:** GPU nodes each contain 4 NVIDIA H100 GPUs, with 94 GiB of HBM2e memory.
Up to 352 NVIDIA H100 GPUs.






Distributed Prover

Witness Computation


How do we distribute the witness computation?



1. **Master** process evaluates WC
2. Master decides a partition
3. Master distributes the traces to the rest of workers

NO PARALLELIZATION + HIGH MEMORY REQUIREMENTS MASTER + COMM OVERHEAD


How do we distribute the witness computation?



1. All processes compute **redundantly** the WC
2. Same partition algorithm is applied
3. Each process keeps only its “owned” instances

NO PARALLELIZATION + HIGH MEMORY REQUIREMENTS ALL PROCS

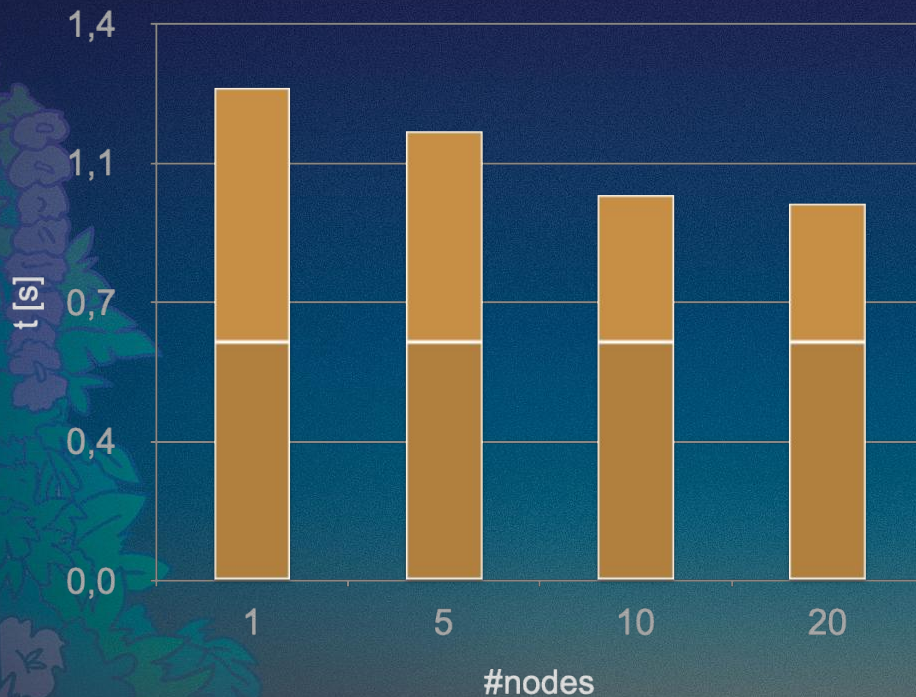
How do we distribute the witness computation?



1. All processes compute a **minimum trace and checkpoints**
2. Same partition algorithm is applied
3. Each process **expands** only its “owned” instances
4. **Reduction** for multiplicity tables

**REDUNDANT MIN TRACE + PARALLEL TRACE EXPANSION +
FINAL REDUCTION**

Results



- Test case: 10K SHA256 (**uncomplete**)
- 40 Instances generated size $N=2^{21}$
 - 21 Main (33 pols)
 - 11 Binary (30 pols)
 - 6 Binary Extended (31 pols)
 - 2 Binary tables (2 pols)
- **Emulator time 0.59s (79 MHz)**
- Traces generator overhead ranges from 110% down to **only 37%**.

Distributed Prover Subproofs

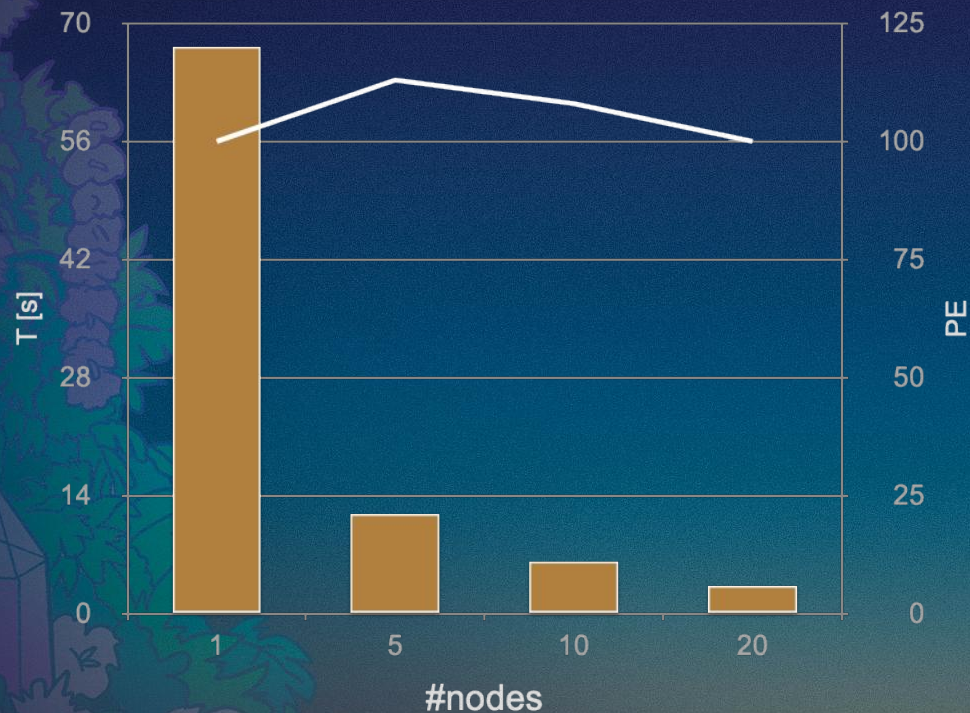
Subproofs generation



1. Redundant transcript (min memory requirements)
2. **All gather** communication to add all roots in each transcript


SYNCHRONIZATION + COMMUNICATION COSTS

Result



- Test case: 10K SHA256 (uncomplete)
- Comparing prover with one node vs distributed prover up to 20 nodes using two process per node
- **100%** Parallel efficiency with 20 nodes
- **Speedup 20x**
- **Min time 3.3s**
- Superlinear with 5 and 10 processes??

Result




- Test case: 10K SHA256 (uncomplete)
- Comparing distributed prover with one node vs distributed prover up to 20 nodes using two process per node
- **80%** Parallel efficiency with 20 nodes
- **Min time 3.3s**
- **Speedup 16x**

Distributed Prover Recursion




Recursion




Point to point comm.

Result



- Test case: 10K SHA256 (uncomplete)
- **Parallel Efficiency 44%**
- Parallelism is limited at the higher levels of aggregation..

Result




- Test case: 10K SHA256 (uncomplete)
- **Parallel Efficiency 88%**
- **Unused resources can be released**




Distributed Prover **Complete**

Result



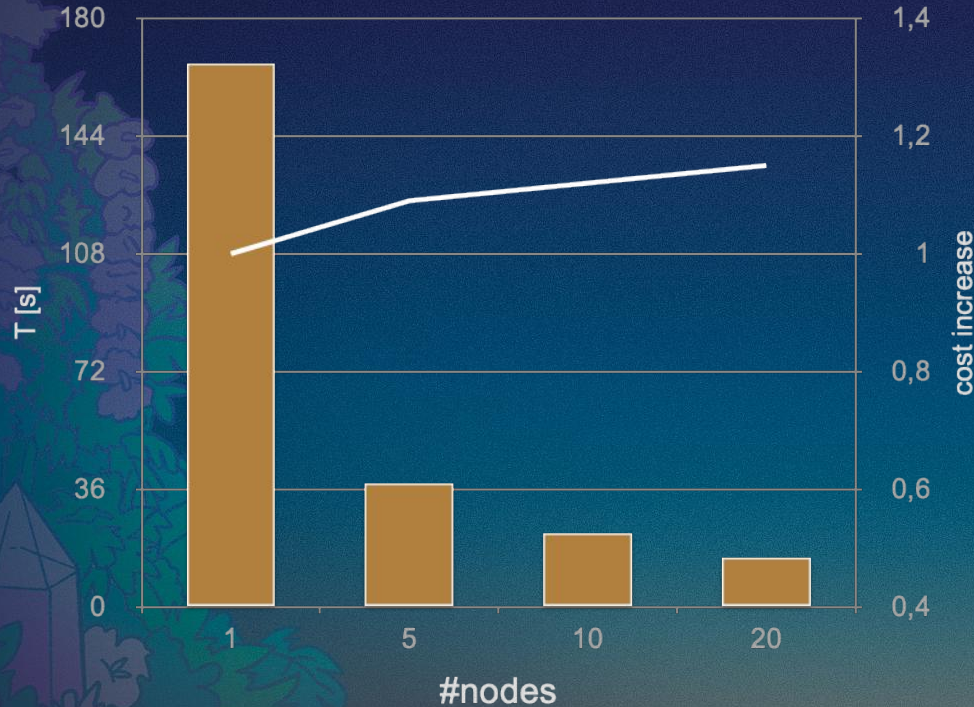
1. Prover time reduced from 125s down to 12.2s
2. **Parallel efficiency of 78%**

Result



1. Prover time reduced from 125s down to 12.2s
2. Parallel efficiency of 78%
3. **Recursion ends up representing 65% of the proof (7.9s)**

Result



1. Prover time reduced from 125s down to 12.2s
2. Parallel efficiency of 78%
3. Recursion ends up representing 65% of the proof (7.9s)
4. **Proof costs increase only 27% while latency proof decreases by 10.3x**



*Distributed computing can reduce latency with
minimal costs increase*

Next steps

- Improve **recursion performance**
 - Algorithmic improvements: STIR, Multi-FRI
 - Improve Circom WC performance
- **GPU** STARK implementation (almost ready)
- Distributed prover on **cloud infrastructure**

Thank you!

Ricard Borrell

Staff Engineer, Polygon
rborrell@polygon.technology
[@Rickb80](https://twitter.com/Rickb80)

