

Optimize zkEVM throughput: Series II

Bangkok 14th November 2024 - Devcon VII

Carlos Matallana
Protocol Team Lead
Ignasi Ramos
Protocol Team Engineer



 **polygon** zkEVM

README

<https://github.com/0xPolygonHermesz/zkevm-rom/blob/devcon-24/README-workshop.md>



Presentation Outline



- 1. zkASM
 - Overview
 - Registers
 - Instructions
 - Examples
 - zkCounters
- 2. zkASM optimizations
- 3. Code: **opPUSH**
 - Diagram
 - Tests
- 4. Code: **mloadX/mstorex**



1

zkASM


zkASM is the language developed by the team that is used to write the program that a compiler will build and the executor will interpret in order to build the execution trace.

```
1 STEP => A
2 0 :ASSERT ; Ensure it is the beginning of the execution
3
4 CTX :MSTORE(forkID)
5 CTX - %FORK_ID :JMPNZ(failAssert)
6
7 B :MSTORE(oldStateRoot)
```

zkASM: Overview

zkASM Compiler

We have implemented a [zkASM compiler](#) that reads a zkASM specification file and compiles it to an output file with the list steps and instructions which the executor will consume in order to compute the execution trace.



Let's check the code!

zkASM: Registers Types

- 8 slots registers → [**V0**, **V1**, ..., **V7**]
 - Each element: Goldilocks number
 - Prime field: $2^{64} - 2^{32} + 1$
 - 63.99 bits
 - 32 bits are used for each slot
 - $8 * 32 = 256$ bits
- 1 slot register → **V0**
 - One element: Goldilocks number
 - Prime field: $2^{64} - 2^{32} + 1$
 - 63.99 bits

zkASM: Registers

- 8 slots registers → [**V0**, **V1**, ..., **V7**]
 - Generic registers: **A**, **B**, **C**, **D**, **E**
 - **SR**: state root
 - **ROTL_C**: Rotate left register **C**
- 1 slot register → **V0**
 - zkCounters: **CNT_ARITH**, **CNT_BINARY**, **CNT_KECCAK_F**, **CNT_MEM_ALIGN**, **CNT_PADDING_PG**, **CNT_POSEIDON_G**, **CNT_STEPS**, **CNT_SHA256**
 - **CTX**: context number
 - **SP**: stack pointer
 - **PC**: smart contract program counter
 - **GAS**: available transaction gas
 - **zkPC**: zkrom program counter
 - **RR**: return register
 - **STEP**: zkrom step counter
 - **MAXMEM**: tracks maximum memory
 - **HASHPOS**: hash position

zkASM: Instructions (0)

- **MLOAD**: load from memory
 - $op = mem[addr]$
- **MSTORE**: save to memory
 - $mem[addr] = op$
- **HASHK**: add bytes to keccak
 - $hashK[E].data[HASHPOS \dots HASHPOS + D - 1] = op[0 \dots D - 1]$
 - **HASHPOS** := **HASHPOS** + **D**
- **HASHKLEN**: digest keccak
 - $hashK[E].length = op$
 - $hashK[E].digest = keccak(hashK[E].data)$
- **HASHKDIGEST**: retrieve keccak hash digest
 - $hashK[E].digest = op$

zkASM: Instructions (1)

- **HASHP**: add bytes to linear poseidon
 - $\text{hashP}[\mathbf{E}].\text{data}[\mathbf{HASHPOS} \dots \mathbf{HASHPOS} + \mathbf{D} - 1] = \text{op}[0 \dots \mathbf{D} - 1]$
 - $\mathbf{HASHPOS} := \mathbf{HASHPOS} + \mathbf{D}$
- **HASHPLEN**: digest linear poseidon
 - $\text{hashP}[\mathbf{E}].\text{length} = \text{op}$
 - $\text{hashP}[\mathbf{E}].\text{digest} = \text{poseidon}(\text{hashP}[\mathbf{E}].\text{data})$
- **HASHPDIGEST**: retrieve linear poseidon hash digest
 - $\text{hashP}[\mathbf{E}].\text{digest} = \text{op}$
- **JMP**: jump to specific zk program counter
 - $\text{zkPC}' = \text{addr}$
- **JMPN**: jump if negative
 - $\text{zkPC}' = (\text{op} < 0) ? \text{addr} : \text{zkPC} + 1;$
- **JMPC**: jump if carry is positive
 - $\text{zkPC}' = (\text{carry} == 1) ? \text{addr} : \text{zkPC} + 1;$

zkASM: Instructions (2)

- **CALL**: jump to specific routine and set return register
 - $zkPC' = addr$
 - **RR** = op0 (**CONST**)
- **RETURN**: go back to return register
 - $zkPC' = addr$ (loaded from RR)
- **ASSERT**: assertion
 - **A** == **op**
- **SLOAD**: storage state-tree load
 - $key0 = [C0, C1, C2, C3, C4, C5, C6, C7]$
 - $key1 = [A0, A1, A2, A3, A4, A5, B0, B1]$
 - $key = HP(key1, HP(key0))$
 - $op = storage.get(SR, key)$

zkASM: Instructions (3)

- **SSTORE**: storage state-tree store
 - $\text{key0} = [\text{C0}, \text{C1}, \text{C2}, \text{C3}, \text{C4}, \text{C5}, \text{C6}, \text{C7}]$
 - $\text{key1} = [\text{A0}, \text{A1}, \text{A2}, \text{A3}, \text{A4}, \text{A5}, \text{B0}, \text{B1}]$
 - $\text{value} = [\text{D0}, \text{D1}, \text{D2}, \text{D3}, \text{D4}, \text{D5}, \text{D6}, \text{D7}]$
 - $\text{SR}' = \text{storage.get}(\text{SR}, \text{key}, \text{value})$
- **ARITH**: arithmetic state machine
 - $\text{A} * \text{B} + \text{C} = \text{D} * (2^{256}) + \text{op}$
- **ADD**: addition
 - $\text{A} + \text{B} = \text{op}$ (set carry as overflow)
- **SUB**: subtraction
 - $\text{A} - \text{B} = \text{op}$ (set carry as underflow)
- **LT**: less than
 - $(\text{A} < \text{B}) = \text{op}$ (set carry as a result)

zkASM: Instructions (4)

- **SLT**: signed less than
 - $\text{signed}(\mathbf{A} < \mathbf{B}) = \text{op}$ (set carry as a result)
- **EQ**: equal
 - $(\mathbf{A} == \mathbf{B}) = \text{op}$ (set carry as a result)
- **AND**: bitwise logic gate **AND**
 - $\mathbf{A} \& \mathbf{B} = \text{op}$
- **OR**: bitwise logic gate **OR**
 - $\mathbf{A} | \mathbf{B} = \text{op}$
- **XOR**: bitwise logic gate **XOR**
 - $\mathbf{A} ^ \mathbf{B} = \text{op}$

zkASM: Instructions (5)

- **MEM_ALIGN_RD**: memory align read
 - MemorySlot0=**A**, MemorySlot1=**B**, Value=op, Offset=**C**
 - $\text{Value} == (\text{MemorySlot0} \ll \text{offset} * 8) \mid (\text{MemorySlot1} \gg 256 - \text{offset} * 8)$
- **MEM_ALIGN_WR**: memory align write 32 bytes
 - MemorySlot0=**A**, MemorySlot1=**B**, Value=op, Offset=**C**, WriteSlot0=**D**, WriteSlot1=**E**
 - $_W0 = (\text{MemorySlot0} \& (2^{256} - 2^{(256 - \text{offset} * 8)})) \mid (\text{Value} \gg \text{offset} * 8)$
 - $_W1 = (\text{MemorySlot1} \& (2^{256} - 1 \gg \text{offset} * 8)) \mid (\text{Value} \ll 256 - \text{offset} * 8)$
 - $_W0 === \text{WriteSlot0}, _W1 === \text{WriteSlot1}$
- **MEM_ALIGN_WR8**: memory align write 8 bit
 - MemorySlot0=**A**, Value=op, Offset=**C**, WriteSlot0=**D**
 - $_W0 = (\text{MemorySlot0} \& (\text{MaskByte} \gg \text{offset} * 8)) \mid ((\text{Value} \& 0xFF) \ll 8 * (31 - \text{offset}))$
 - $_W0 === \text{WriteSlot0}$

Generic registers

A, B, C, D, E -> 8 elements of goldilock prime field number -> 32 bits -> 256 bits

A = A7, A6, A5, A4, A3, A2, A1, A0

Registry assignments:

- $5 \Rightarrow A \text{ // } A7-A1 = [0], A0 = [5]$
- $A \Rightarrow C \text{ // } A == C$
- $4294967295 (2^{32} - 1) \Rightarrow B \text{ // } B7-B1 = [0], B0 = [0xffffffff]$
- $A + B \Rightarrow B \text{ // } B7-B1 = [0], B0 = [0x100000004] \text{ ! OVERFLOW}$

ROTL_C

C = 0x 9146dd22 7e54eb92 21ee4527 2735a8ae 52de694a f5749e7a b473b3f1
1328c778
; [7], [6], [5], [4], [3], [2], [1], [0]

ROTL_C = 0x 7e54eb92 21ee4527 2735a8ae 52de694a f5749e7a b473b3f1
1328c778 9146dd22
; [6], [5], [4], [3], [2], [1], [0], [7]

JMPN - JMPZ - JMPNZ

JMPN: jump if op[0] is negative

A :JMPN(jumplfNeg) // jump to *jumplfNeg* if A[0] is neg (more than 32 bits)

JMPZ: jump if op[0] was zero

A :JMPZ(jumplfZero) // jump to *jumplfZero* if A[0] is zero

JMPNZ: jump if op[0] was different of zero

A :JMPNZ(jumplfNoZero) // jump to *jumplfNoZero* if A[0] is not zero


JMPC: jump if carry bit, only use with binary operations

JMPNC: jump if no carry bit, only use with binary operations

\$:LT, JMPC(carry, noCarry) // jump to *carry* if A is LT B, else jump no *noCarry*

HASHK - HASHKLEN - HASHKDIGEST

- **HASHK**: add bytes to linear poseidon. Dependant on **HASHPOS**, **D** registers
 - **HASHPOS** is the offset, **D** is the size
 - $\text{hashK}[\mathbf{E}].\text{data}[\mathbf{HASHPOS} \dots \mathbf{HASHPOS} + \mathbf{D} - 1] = \text{op}[0 \dots \mathbf{D} - 1]$
 - **HASHPOS** := **HASHPOS** + **D**
- **HASHKLEN**: close hash with length **HASHPOS**
- **HASHKDIGEST**: get the digest



```
1777 32 => HASHPOS
1778 32 => D
1779 A :HASHK(E)
```

ASSERT

- Breaks verification if
 - **A** **!=** **op**
 - Code example:

1656	1 => A
1657	2 :ASSERT

Free Input

- Loads any value into a register
 - Value is not verified. It is freely chosen by the prover
 - Support basic syntax operations: **+**, **-**, *****, **&**, **||**, ...

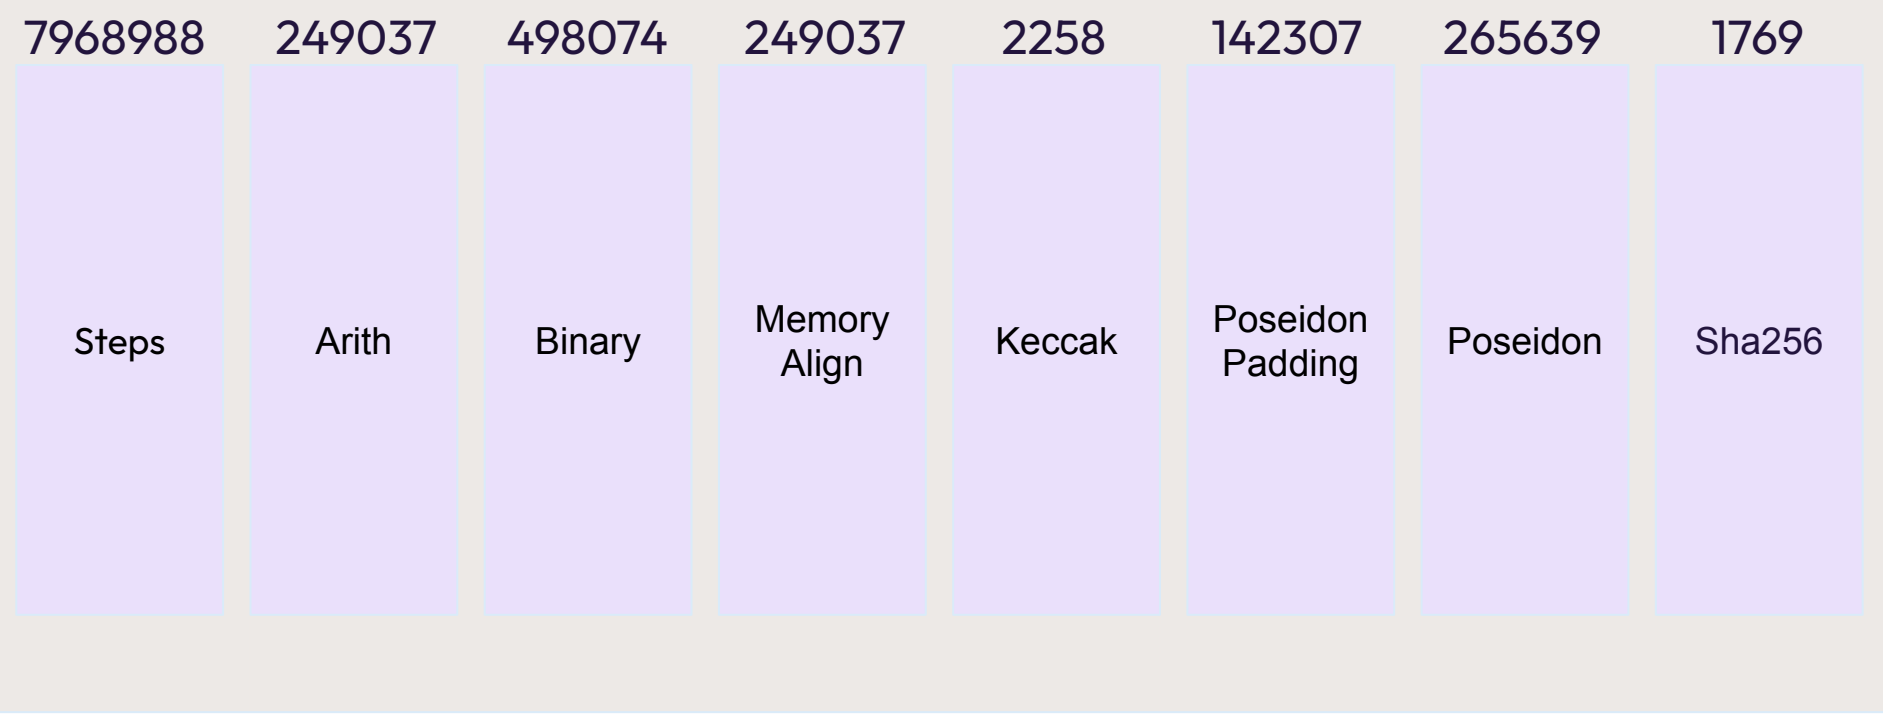
1813	#{B & 0x00000001}	:JMPZ(hashLeft0)
------	-------------------	------------------

zkASM: zkCounters

```
104 ; COUNTERS
105 CONST %MIN_STEPS_FINISH_BATCH = 200 ; min steps to finish tx
106 CONST %TOTAL_STEPS_LIMIT = 2**23
107
108 CONST %MAX_CNT_STEPS_LIMIT = %TOTAL_STEPS_LIMIT - %MIN_STEPS_FINISH_BATCH
109 CONST %MAX_CNT_ARITH_LIMIT = %TOTAL_STEPS_LIMIT / 32
110 CONST %MAX_CNT_BINARY_LIMIT = %TOTAL_STEPS_LIMIT / 16
111 CONST %MAX_CNT_MEM_ALIGN_LIMIT = %TOTAL_STEPS_LIMIT / 32
112 CONST %MAX_CNT_KECCAK_F_LIMIT = (%TOTAL_STEPS_LIMIT / 155286) * 44
113 CONST %MAX_CNT_PADDING_PG_LIMIT = (%TOTAL_STEPS_LIMIT / 56)
114 CONST %MAX_CNT_POSEIDON_G_LIMIT = ((%TOTAL_STEPS_LIMIT / 30))
115 CONST %MAX_CNT_SHA256_F_LIMIT = ((%TOTAL_STEPS_LIMIT - 1) / 31488) * 7
116
117 CONST %SAFE_RANGE = 20 ; safe guard counters to not take into account (%RANGE = 1 / SAFE_RANGE)
118
119 CONST %MAX_CNT_STEPS = %MAX_CNT_STEPS_LIMIT - (%MAX_CNT_STEPS_LIMIT / %SAFE_RANGE)
120 CONST %MAX_CNT_ARITH = %MAX_CNT_ARITH_LIMIT - (%MAX_CNT_ARITH_LIMIT / %SAFE_RANGE)
121 CONST %MAX_CNT_BINARY = %MAX_CNT_BINARY_LIMIT - (%MAX_CNT_BINARY_LIMIT / %SAFE_RANGE)
122 CONST %MAX_CNT_MEM_ALIGN = %MAX_CNT_MEM_ALIGN_LIMIT - (%MAX_CNT_MEM_ALIGN_LIMIT / %SAFE_RANGE)
123 CONST %MAX_CNT_KECCAK_F = %MAX_CNT_KECCAK_F_LIMIT - (%MAX_CNT_KECCAK_F_LIMIT / %SAFE_RANGE)
124 CONST %MAX_CNT_PADDING_PG = %MAX_CNT_PADDING_PG_LIMIT - (%MAX_CNT_PADDING_PG_LIMIT / %SAFE_RANGE)
125 CONST %MAX_CNT_POSEIDON_G = %MAX_CNT_POSEIDON_G_LIMIT - (%MAX_CNT_POSEIDON_G_LIMIT / %SAFE_RANGE)
126 CONST %MAX_CNT_SHA256_F = %MAX_CNT_SHA256_F_LIMIT - (%MAX_CNT_SHA256_F_LIMIT / %SAFE_RANGE)
127 CONST %MAX_CNT_POSEIDON_SLOAD_SSTORE = 518
```

zkASM: zkCounters

Batch





2

zkASM optimizations

zkASM Optimizations: **SAVE, RESTORE**

- Feature to save and restore registers. The registers involved are B, C, D, E, RR, RCX, and RID. Also, the **op** was saved
- Useful when a routine uses instructions that must use some registers, so it is needed to save/recover those registers. Save steps

```
591  VAR GLOBAL tmpB
592  VAR GLOBAL tmpC
593  VAR GLOBAL tmpD
594  VAR GLOBAL tmpE
595  function:
596  B          :MSTORE(tmpB)
597  C          :MSTORE(tmpC)
598  D          :MSTORE(tmpD)
599  E          :MSTORE(tmpE)
600  ;; do some function
601  $ => B     :MLOAD(tmpB)
602  $ => C     :MLOAD(tmpC)
603  $ => D     :MLOAD(tmpD)
604  $ => E     :MLOAD(tmpE)
605           :RETURN
```



```
607  function:
608  |          |          :SAVE(B, C, D, E, RR, RCX)
609  ;; do some function
610  |          |          :RESTORE
611  |          |          :RETURN
```

zkASM Optimizations: Assume Free Input

With the internal **assumeFree** flag, we indicate to the processor that it has to verify the FREE INPUT, not the **op**. This behavior could be used with memory-style instructions like **MLOAD** and **HASHx**. In assembly, to indicate it, we use the prefix **F_** on instructions, like **F_MLOAD**, **F_HASHP**, etc. For example:

```
; increment value of B with contents of var1

; using MLOAD (worse case) - 4 steps
C           :MSTORE(tmpC)
$ => C      :MLOAD(var1)
B + C => B
$ => C      :MLOAD(tmpC)

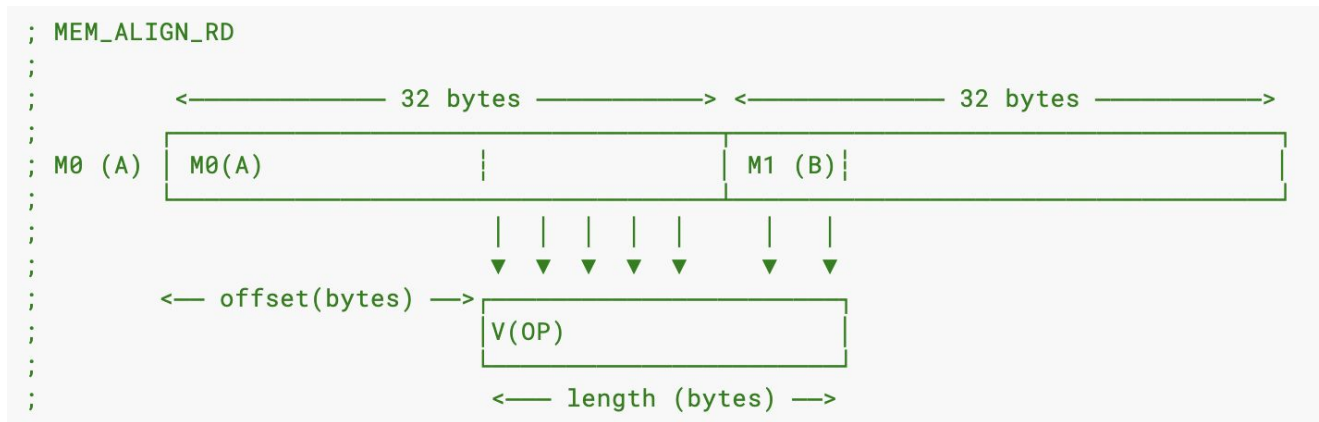
; using F_MLOAD - 1 step
$ + B => B  :F_MLOAD(var1)
```


zkASM Optimizations: MEM_ALIGN with variable length

- One previous limitation of MEM_ALIGN was that it only worked with 32-byte values, and it wasn't useful to read bytes over 256-bit memory.
- These extra features are right/left alignment and little/big (by default right-alignment and big-endian).
- To avoid using more registers, all these parameters are “coded”
 - $\text{mode} = \text{offset_bytes} (0-64)$
 - $+ 128 * \text{length} (0-32, 0 \text{ for compatibility is equivalent to } 32)$
 - $+ 8192 * \text{left_alignment} (0-1)$
 - $+ 16384 * \text{little_endian} (0-1)$


zkASM Optimizations: MEM_ALIGN with variable length

To read from memory, we need to read two consecutive addresses (M0, M1) and store them on registers A and B, respectively. Use C to store mode and OP to obtain bytes read.



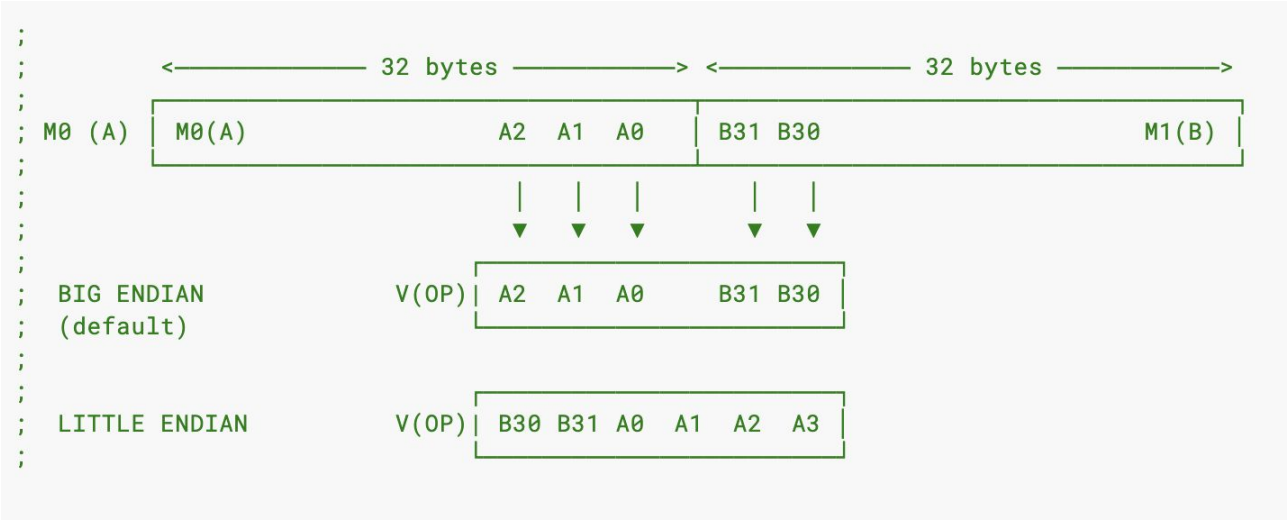
zkASM Optimizations: MEM_ALIGN with variable length

Another parameter of mem_align is its alignment, which defines how bytes read are stored inside the 32-byte register. If the length is 32, the left and right alignments are the same.



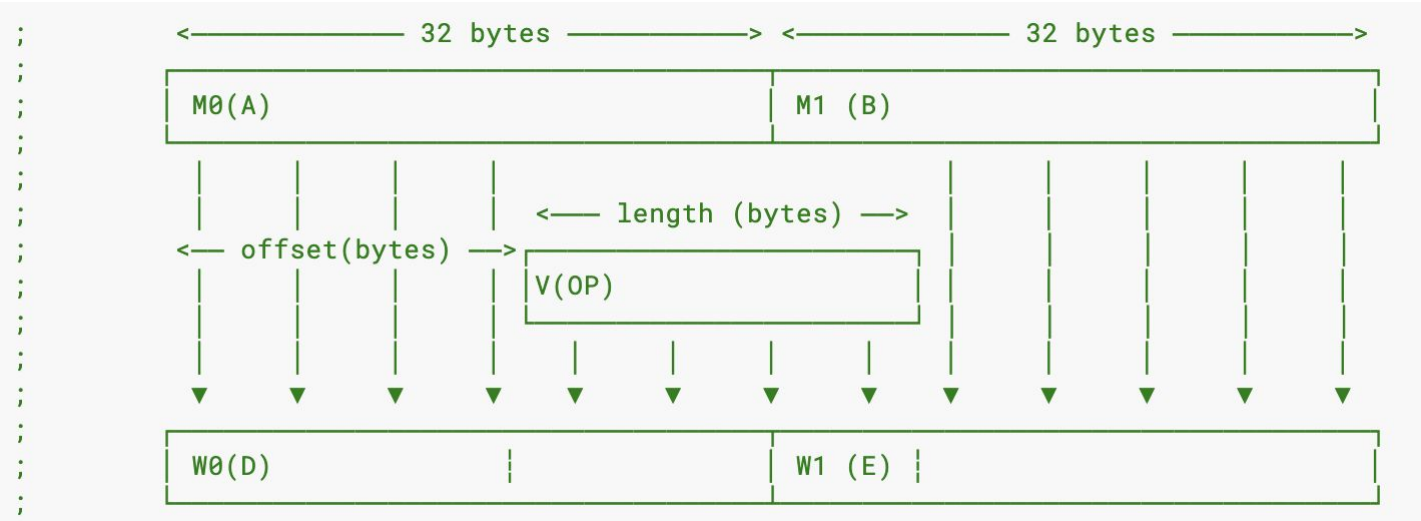
zkASM Optimizations: MEM_ALIGN with variable length

The last parameter of mem_align is its endian type, which defines how the ordered bytes are read. Big endian takes bytes as they are in memory, while little endian reverses them.



zkASM Optimizations: MEM_ALIGN with variable length

Finally, operation MEM_ALIGN_WR, in this case, verifies D as W0 and E as W1, indicating how memory M0 and M1 were modified after "writing" n bytes of OP with offset. The extra parameters alignment and endian define how to take this byte from V (OP).



zkASM Optimizations: MEM_ALIGN with variable length

```
; #3168 w=[101376-101407]
; -----.-----.---#####
0x101112131415161718191A1B1C1D1E1F202122232425262728292A2B2C2D2E2Fn => A
; #####-----.-----.-----
0x303132333435363738393A3B3C3D3E3F404142434445464748494A4B4C4D4E4Fn => B
12n => C
$ => A :MEM_ALIGN_RD
0x1C1D1E1F202122232425262728292A2B2C2D2E2F303132333435363738393A3Bn :ASSERT
```

```

; #3172 w=[101504-101535]
0x101112131415161718191A1B1C1D1E1F202122232425262728292A2B2C2D2E2Fn => A
; -----#####
0x101112131415161718191A1B0E1E2E3E4E5E6E7E8E9EAECEDEEEEF0F1F2F3Fn => D
12n => C
0x303132333435363738393A3B3C3D3E3F404142434445464748494A4B4C4D4E4Fn => B
; #####-----
0x4F5F6F7F8F9FAFBFCFDFEFFF3C3D3E3F404142434445464748494A4B4C4D4E4Fn => E
0x0E1E2E3E4E5E6E7E8E9EAECEDEEEEF0F1F2F3F4F5F6F7F8F9FAFBFCFDFEFFFn :MEM_ALIGN_WR

```

zkASM Optimizations: ARITH_MOD (modular arithmetic)

- Given registers **A,B,C,D,op** it checks whether the linear operation **A*B + C** is equal to **op** modulo **D**
- Useful when doing modulo operations (ex: ecrecover, mulmod)

$$A \cdot B + C = \text{op} \pmod{D}$$

; Compute A%D and return it on register C

checkRem:

; A · 1 + 0 = op (mod D)

1 => B

0 => C

\$(A % D) => E :ARITH_MOD

E => C :RETURN



3

Code: opPUSH



opPUSH: Diagram

- Bytecode is loaded at the very beginning of the transaction and it must match the **hashBytecode** in the state-tree



```
429      ; get hash contract
430      $ => A                               :MLOAD(txDestAddr)
431      %SMT_KEY_SC_CODE => B
432      $ => A                               :SLOAD
433
434      ; get a new hashPID
435      $ => E                               :MLOAD(nextHashPID)
436      E                                   :MSTORE(contractHashId)
437      E+1                               :MSTORE(nextHashPID)
438
439      ; load contract bytecode
440      A                                   :HASHPDIGEST(E)
```

- Limitations:
 - Bytes can be read as many times as it needs but it must be always in the same way (offset & size)
 - Always reads bytes 1 by 1

opPUSH: Diagram



1 byte




opPUSH: Diagram

5 byte

Final
Value:



Bytecode:



opPUSH: Tests

- Load an array of bytes as if it is a bytecode:

```
17      ; add bytes one by one
18      0x0102030405060708090A0B0C0D0E0F101112131415161718191A1B1C1D1E1F20n => A      :CALL(initLoop)
19      0x2122232425262728292A2B2C2D2E2F303132333435363738393A3B3C3D3E3F40n => A      :CALL(initLoop)
20      0x4142434445464748494A4B4C4D4E4F505152535455565758595A5B5C5D5E5F60n => A      :CALL(initLoop)
```



0x010203... 5D5E5F60

- Perform multiple PUSH opcodes with variable length:

```
36      ;; Start test PUSH1
37      ; PUSH1 reading 1 byte at position 0 in the bytecode
38      0 => PC
39      1 => D,E                      :CALL(readPush)
40      0x01 => A
41      E                          :ASSERT
```

opPUSH: Tests

0x010203...

...

...

...

...

...

...

...5D5E5F60


```
80      ;; Start test PUSH3
81      ; PUSH3 reading 3 byte at position 0 in the bytecode
82      0 => PC
83      3 => D,E                      :CALL(readPush)
84      0x010203 => A
85      E                          :ASSERT
```

```
283      ;; Start test PUSH32
284      ; PUSH32 reading 32 byte at position 42 in the bytecode
285      42 => PC
286      32 => D,E                      :CALL(readPush)
287      0x2B2C2D2E2F303132333435363738393A3B3C3D3E3F404142434445464748494An => A
288      E                          :ASSERT
289
290      0 => A,B,C,D,E,CTX, SP, PC, GAS, SR, HASHPOS, RR ; Set all registers to 0
291      :JMP(finalizeExecution)
```

Performance comparison

OLD

```
{  
  cntArith: 0n,  
  cntBinary: 278n,  
  cntKeccakF: 0n,  
  cntSha256F: 0n,  
  cntMemAlign: 0n,  
  cntPoseidonG: 2n,  
  cntPaddingPG: 2n,  
  cntSteps: 6417  
}
```



NEW

```
{  
  cntArith: 0n,  
  cntBinary: 1n,  
  cntKeccakF: 0n,  
  cntSha256F: 0n,  
  cntMemAlign: 0n,  
  cntPoseidonG: 2n,  
  cntPaddingPG: 2n,  
  cntSteps: 1068  
}
```



3

Code: mloadX/mstoreX

Example before optimization:

bytesToStore:

```
0x00000000FF00000000000000000000000000000000000000000000000000000
```

Offset = 4 bytes

Current memory:

0xAAAAAAAA**A**AAA

1- Shift left (32 - offset - length) + shift right

0x00000000**00**AAA

2- Shift right (32 - offset - length) + shift left

[illegible]

3- SUM

[illegible]

4- MemAlign

Before optimization

cntArith: 7
cntBinary: 12
cntKeccakF: 0
cntSha256F: 0
cntMemAlign: 0
cntPoseidonG: 0
cntPaddingPG: 0
cntSteps: 240



After optimization

cntArith: 0
cntBinary: 0
cntKeccakF: 0
cntSha256F: 0
cntMemAlign: 1
cntPoseidonG: 0
cntPaddingPG: 0
cntSteps: 25

Thank you!

Carlos Matallana

Protocol and Integration Team, Polygon zkEVM
carlos@polygon.technology



@KrlsMata

Ignasi Ramos

Protocol and Integration Team, Polygon zkEVM
ignasi@polygon.technology



@0xIgnasi

