

# Security of Fiat-Shamir transformation

Michał Zając  
Nethermind

# Zero-knowledge proofs



$x$  — statement, e.g. “Circuit  $C$  executed on input  $a$  outputs  $b$ ”

$w$  — witness, e.g., some private inputs to the program  $P$ , intermediate values of  $C$ ’s evaluation

**Completeness** Honest  $V$  always accepts a proof from a prover  $P$

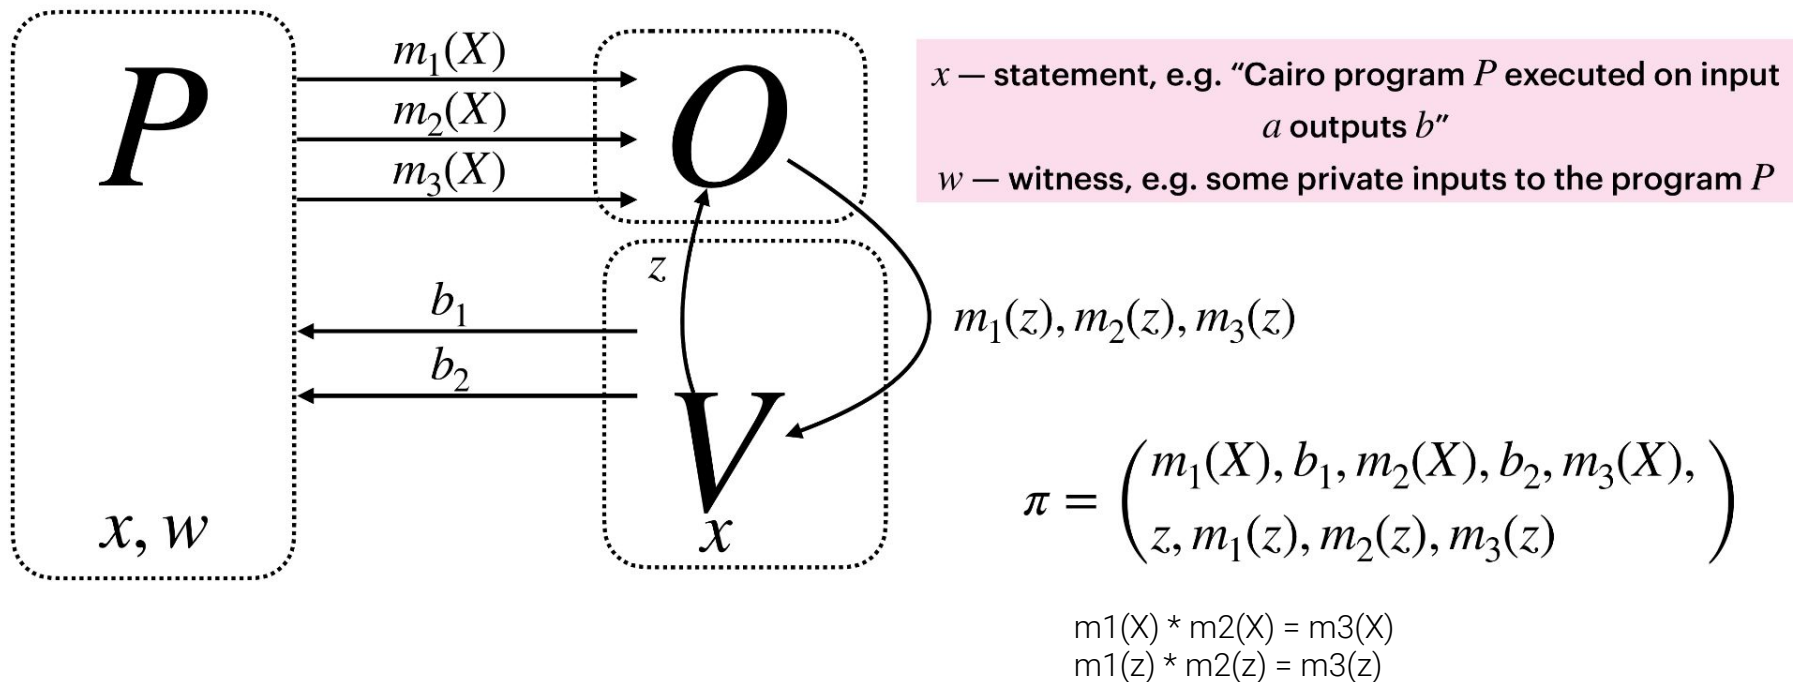
**Soundness** Probability that  $P^*$  makes  $V$  accept false statement is negligible

**Knowledge soundness** If  $P^*$ ’s proof is accepted by  $V$ , then  $P$  knows  $w$

**Zero-knowledge**  $V$  learns nothing new about  $x$ , except that  $Rel(x, w)$

**How SNARKs are built?**

# From PIOP to SNARKs




**Problem 1** this proof is huge and we rely on  $O$


# Polynomial commitments

Allow the **prover** to send a **short digest** that represents the polynomial **without revealing it**

Allow the **verifier** to **verifiably evaluate** the polynomial represented by the digest



# From PIOP to SNARKs. Replace $\mathcal{O}$




$c_i$  polynomial commitment to polynomial  $m_i(X)$   
 $\pi_i$  proof for the polynomial commitment

$$\pi = \begin{pmatrix} c_1, b_1, c_2, b_2, c_3, \\ z, m_1(z), m_2(z), m_3(z) \\ \pi_1, \pi_2, \pi_3 \end{pmatrix}$$

**Problem 2** the protocol is interactive

# From PIOP to SNARKs. Fiat—Shamir transformation



$c_i$  polynomial commitment to polynomial  $m_i(X)$

$\pi_i$  proof for the polynomial commitment

$R$  random oracle

$$tr_1 = x, c_1 \qquad b_1 = R(tr_1)$$


$$tr_2 = tr_1, b_1, c_2 \qquad b_2 = R(tr_2)$$

$$tr_3 = tr_2, c_2, b_2 \qquad z = R(tr_3)$$

$$\pi = \begin{pmatrix} c_1, b_1, c_2, b_2, c_3, \\ z, m_1(z), m_2(z), m_3(z) \\ \pi_1, \pi_2, \pi_3 \end{pmatrix}$$

**Problem 3** random oracle doesn't exist

# From PIOP to SNARKs. Removing random oracle



$c_i$  polynomial commitment to polynomial  $m_i(X)$

$\pi_i$  proof for the polynomial commitment

$H$  hash function

$$tr_1 = x, c_1 \quad b_1 = h(tr_1)$$

$$tr_2 = tr_1, b_1, c_2 \quad b_2 = h(tr_2)$$

$$tr_3 = tr_2, c_2, b_2 \quad z = h(tr_3)$$

$$\pi = \begin{pmatrix} c_1, b_1, c_2, b_2, c_3, \\ z, m_1(z), m_2(z), m_3(z) \\ \pi_1, \pi_2, \pi_3 \end{pmatrix}$$

**Problem 4** how secure is such protocol?




# Fiat—Shamir transformation security

# The power of rewinding

## Intuition

**interactive:** the attacker must give only good answers

**non-interactive:** the attacker can cheat by finding a challenge from a lucky set



**Lucky set** Set of challenges  $B$  such that if  $P^*$  gets  $b \in B$  as a challenge then  $V$  accepts and  $\Pi$  enters *lucky state*

$V$  accepts **only if** the protocol is in the *lucky state*

$\epsilon$  — probability that  $P^*$  makes  $V$  accept a false statement

$\eta$  — the probability that  $\Pi$  gets into the lucky state at  $i$ -th round

$(1 - \eta)^k$  — the probability that  $\Pi$  doesn't get into the lucky state in one of the  $k$  rounds


$$\epsilon = 1 - (1 - \eta)^k \approx k \cdot \eta$$

# How to compare the security?

$\epsilon$  - probability that  $P^*$  convinces  $V$  on a false statement (a.k.a.  $P^*$  breaks the soundness)

$\eta$  - probability that  $P^*$  gets a challenge from a lucky set

$Q$  - number of hashes  $P_{fs}^*$  can make (a.k.a.  $P_{fs}^*$ 's running time or *attack budget*)



$1 - \eta$ : probability that  $P^*$  **doesn't get a challenge from a lucky set** in a particular round

$(1 - \eta)^{Q/k}$ : probability that  $P^*$  doesn't get a challenge from a lucky set in a particular round **in  $Q/k$  trials**

$\prod_{i=1}^k (1 - \eta)^{Q/k}$ : probability that  $P^*$  doesn't get a challenge from a lucky set in a particular round in  $Q/k$  trials **for each of  $k$  rounds**

$$\epsilon_{fs} = 1 - \prod_{i=1}^k (1 - \eta)^{Q/k} = 1 - (1 - \eta)^Q \approx \eta \cdot Q$$

# But what's the security?

$$\epsilon \approx 2^{-100}$$

What's  $Q$ ?

$$600M * 10^{12} \approx 6 * 2^{69} [h/s]$$
$$2^{85} [h/day]$$

$$\epsilon' \approx 2^{-100} \cdot Q = 2^{-15}$$


## Total Hash Rate (TH/s)

The estimated number of terahashes per second the bitcoin network is performing in the last 24 hours.



# Bootstrapping security with parallel repetition

# Parallel repetition




$\epsilon$  probability of breaking the security of  $\Pi$

$\epsilon_2$  - probability of breaking  $\Pi^2$

$$\epsilon_2 = \epsilon \cdot \epsilon$$

# Parallel repetition. FS transformation



transcripts of  $\Pi_{fs,1}, \Pi_{fs,2}$  depend on each other

$\epsilon_{fs}$  probability of breaking the security of  $\Pi_{fs}$

$\epsilon_{fs,2}$  - probability of breaking  $\Pi_{fs}^2$

$$\epsilon_{fs,2} \neq \epsilon_{fs} \cdot \epsilon_{fs}$$

# Lucky sets in FS transformation

**Lucky set** Set of challenges  $B$  such that if  $P^*$  gets  $b \in B$  as a challenge then  $V$  accepts


$B_i^j$  — lucky set for  $i$ -th challenge in execution  $j$

$N$  — set of all challenges

$\eta$  — max probability that  $\Pi_i$  enters *lucky state* in  $B_1^j$

$k$  — number of protocol rounds

$t$  — number of executions



$P^*$  needs to enter a *lucky state* in **all** executions to be successful

$$(1 - (1 - \eta)^k) \cdot \dots \cdot (1 - (1 - \eta)^k) \approx (\eta \cdot k)^t$$

## Example

$$t = 8, k = 4,$$

$$\eta = 2^{-20}$$

$$(2^{-20} \cdot 4)^8 = (2^{-18})^8 = 2^{-144}$$



# Lucky sets in FS transformation

**Lucky set** Set of challenges  $B$  such that if  $P^*$  gets  $b \in B$  as a challenge then  $V$  accepts

$B_i^j$  — lucky set for  $i$ -th challenge in execution  $j$


$N$  — set of all challenges

$\eta_i^j$  — max probability that  $\Pi_i$  enters *lucky state* in

$k$  — number of protocol rounds

$t$  — number of executions

$Q$  —  $P^*$ 's running time



## $P^*$ 's strategy

- (1) Send the first message
- (2) Get the challenge  $b_1^1$
- (3) If  $b_1^1 \notin B_1^1$ , rewind to (1), propose a new  $tr_1^1, tr_1^2$ , get a new challenge  $b_1^1$
- (4) Repeat until  $b_1^1 \in B_1^1$  found
- (5) Send  $tr_2^2$
- (6) Get a challenge  $b_2^2$
- (7) If  $b_2^2 \notin B_2^2$ , rewind to (5), propose new  $tr_2^2$  and get a new challenge  $b_2^2$
- (8) Repeat until  $b_2^2 \in B_2^2$

# Lucky sets in FS transformation

**Lucky set** Set of challenges  $B$  such that if  $P^*$  gets  $b \in B$  as a challenge then  $V$  accepts


$B_i^j$  — lucky set for  $i$ -th challenge in execution  $j$

$N$  — set of all challenges

$\eta_i^j$  — max probability that  $\Pi_i$  enters *lucky state* in

$k$  — number of protocol rounds

$t$  — number of executions



$P^*$  needs to enter a *lucky state* in **all** executions to be successful

$$(1 - (1 - \eta)^{Q/t}) \cdot \dots \cdot (1 - (1 - \eta)^{Q/t}) \approx (Q/t \cdot \eta)^t$$

## Example

$$t = 10, k = 4, Q = 2^{20}$$

$$\eta = 2^{-20}$$

$$(2^{20}/8 \cdot 2^{-20})^8 = (1/8)^{10} = 2^{-30}$$

**There is more...**


# Coordinated disclosure of vulnerabilities

*By Jim Miller*

We've dubbed this class of vulnerabilities Frozen Heart. The word *frozen* is an acronym for FoRging OF ZeRo kNoWledge proofs, and the Fiat-Shamir transformation is at the *heart* of most proof systems: it's vital for their practical use, and it's generally located centrally in protocols. We hope that a catchy moniker will help raise awareness of these issues in the cryptography and wider technology communities.

**TL;DR** Some implementations of Fiat—Shamir transformation didn't include all the necessary information in the hash function's input.

This is a coordinated disclosure: we have notified those parties who we believe have been or may be affected by the issues prior to our public release of the information. The parties who were affected:



Hash needs to be computed on **all the data** that the prover sent so far

# Some more attacks

zkSNARKs & The Last Challenge

## TL;DR don't deviate from the protocol description

*by Oana Ciobotaru, Maxim Peter and Vesselin Velichkov*

[OpenZeppelin](#) recently identified a critical vulnerability during an audit of [Linea](#)'s PLONK verifier. At its core, the issue is similar in nature to some previously disclosed vulnerabilities (e.g., [Frozen Heart](#), [00](#)). The 'Last Challenge' vulnerability arises from the ability of a malicious prover to exploit the degrees of freedom introduced by an incorrect application of the Fiat-Shamir transform when computing the final PLONK challenge. A malicious prover exploiting this vulnerability could steal all the assets in the rollup by submitting a proof for an invalid state transition. While the issue was promptly communicated and [fixed](#), we believe the specifics are worth sharing more broadly so that others may learn to recognize similar patterns and protect themselves accordingly.

Prior to describing the issue in detail below, we give a brief high-level introduction to Zero-Knowledge Succinct Non-interactive Arguments of Knowledge (zkSNARKs) and the Fiat-Shamir transform.

**Be careful when you fine-tune your parameters**